Fruit Market · Decision register

Generated by `npm run records` from `docs/workflow/records/`. Do not edit by hand. Last recorded change: 2026-10-26 23:18 (Asia/Kuala_Lumpur).

Approval status and delivery status are separate. Only the owner's explicit words set "approved".

Pending approvals

ID	What	Status
DEC-001	Adopt the personal workflow guide 1.2 for Fruit Market	Proposed
DEC-004	PDF exports with pandoc and the existing Playwright browser	Proposed
DEC-005	Budget interpretation: RM500 per month across all games	Proposed
DEC-014	Use the Concept and Theme Companion for creative work	Proposed
DEC-020	UI and UX polish in three phases	Awaiting approval
DEC-021	Art direction for the polish: explore A/B/C, recommend toy-box tabletop	Awaiting approval
PLAN-003	Package C: online, capacity and release evidence (outline only)	Proposed
PLAN-007	UI polish phase 1: fix what looks broken	Awaiting approval
PLAN-008	UI polish: art-direction style test and layout mockups (design only)	Awaiting approval

Current decisions

ID	Decision	Approval	Delivery	Reports
DEC-001	Adopt the personal workflow guide 1.2 for Fruit Market	Proposed	In progress	REP-001
DEC-002	Package A scope = PLAN-001	Approved (owner, 2026-10-26)	Verified	REP-001, REP-002
DEC-003	Live deploys only through the manual Deploy workflow after verify	Approved (owner, 2026-10-26)	Verified	REP-001, REP-002
DEC-004	PDF exports with pandoc and the existing Playwright browser	Proposed	Implemented	REP-001
DEC-005	Budget interpretation: RM500 per month across all games	Proposed	Not started	REP-001
DEC-006	Preview = separate staging Worker	Approved (owner, 2026-10-26)	Verified	REP-002, REP-003
DEC-007	Package B scope = PLAN-002	Approved (owner, 2026-10-26)	Verified	REP-003, REP-004
DEC-008	Canonical record format: JSON records plus Markdown narrative reports	Approved (owner, 2026-10-26)	Verified	REP-003, REP-004
DEC-009	Records check runs in verify and CI	Approved (owner, 2026-10-26)	Verified	REP-003, REP-004

ID	Decision	Approval	Delivery	Reports
DEC-010	Reconnect fix scope = PLAN-004 (evidence first, then fix)	Approved (owner, 2026-10-26)	Verified	REP-005, REP-006
DEC-011	Stress test runs on demand, not on every push	Approved (owner, 2026-10-26)	Verified	REP-005, REP-006
DEC-012	4-player clarity scope = PLAN-005 (screen text only)	Approved (owner, 2026-10-26)	Verified	REP-007, REP-008
DEC-013	Publish version 1.1.1 to players	Approved (owner, 2026-10-26)	Not started	REP-009, REP-013
DEC-014	Use the Concept and Theme Companion for creative work	Proposed	In progress	REP-010
DEC-015	Adopt three design pillars	Approved (owner, 2026-10-26)	Implemented	REP-010, REP-011, REP-013
DEC-016	First creative test: reveal-comprehension playtest	Approved (owner, 2026-10-26)	Not started	REP-010, REP-011, REP-013
DEC-017	Setting: a market where the whole world shops	Approved (owner, 2026-10-26)	Implemented	REP-011
DEC-018	Rename Panda Lim while keeping the cast coherent	Approved (owner, 2026-10-26)	Verified	REP-011, REP-012
DEC-019	Rename scope = PLAN-006, shipped in 1.1.2	Approved (owner, 2026-10-26)	Verified	REP-011, REP-012
DEC-020	UI and UX polish in three phases	Awaiting approval	Not started	REP-014, REP-015
DEC-021	Art direction for the polish: explore A/B/C, recommend toy-box tabletop	Awaiting approval	Not started	REP-015

Work packages

ID	Package	Approval	Delivery	Release
PLAN-001	Package A: baseline, one verification path, preview and owner brief	Approved (owner, 2026-10-26)	Verified	1.1.1 candidate
PLAN-002	Package B: canonical records, generated views and checks	Approved (owner, 2026-10-26)	Verified	1.1.1 candidate
PLAN-003	Package C: online, capacity and release evidence (outline only)	Proposed	Not started	Unreleased
PLAN-004	Reconnect fix: the game screen sometimes doesn't return after rejoining an online room	Approved (owner, 2026-10-26)	Verified	1.1.1 candidate
PLAN-005	4-player clarity: make it obvious that up to 4 people can play	Approved (owner, 2026-10-26)	Verified	1.1.1 candidate
PLAN-006	Rename the panda rival (cosmetic, within the animal-trader cast)	Approved (owner, 2026-10-26)	Verified	Unreleased
PLAN-007	UI polish phase 1: fix what looks broken	Awaiting approval	Not started	Unreleased

ID	Package	Approval	Delivery	Release
PLAN-008	UI polish: art-direction style test and layout mockups (design only)	Awaiting approval	Not started	Unreleased

Decision details

DEC-001 · Adopt the personal workflow guide 1.2 for Fruit Market

- **Question:** Should Fruit Market be run with the workflow guide (one approved package at a time, evidence-based checks, decision records)?
- **Choice:** Yes; one package at a time, starting with Package A.
- **Rationale:** The game is complete but releases were unreviewed and records lived in chat history.
- **Player impact:** None directly; fewer broken releases.
- **Technical impact:** Adds checks, records and a release gate.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Proposed
- **Delivery:** In progress
- **Review trigger:** If the process costs more time than it saves after three packages.
- **Reports:** REP-001

DEC-002 · Package A scope = PLAN-001

- **Question:** What should the first setup package contain?
- **Choice:** One verify command, preview before publishing, short records and an owner brief (PLAN-001).
- **Rationale:** Smallest package that stops unreviewed releases and gives the owner a clear way to try builds.
- **Player impact:** None until a release; protects players from silent breakage.
- **Technical impact:** package.json scripts, CI and deploy workflows, new docs.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-001 only, including its listed tests and documentation” (2026-10-26)
- **Delivery:** Verified (PLAN-001)
- **Review trigger:** Acceptance A4 (owner preview playtest) still pending.
- **Reports:** REP-001, REP-002

DEC-003 · Live deploys only through the manual Deploy workflow after verify

- **Question:** When should the live game change?
- **Choice:** Only when the owner runs Deploy LIVE game and types PUBLISH; it runs verify first and a smoke test after. Never on push.
- **Rationale:** Makes 'a preview is not a release' true for this repository.

- **Player impact:** Updates arrive only after the owner approves them.
- **Technical impact:** deploy-web.yml is manual only.
- **Data and security:** None.
- **Cost:** RM0; one extra tap per release.
- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-001 only, including its listed tests and documentation” (2026-10-26)
- **Delivery:** Verified (PLAN-001)
- **Review trigger:** If urgent fixes are delayed because nobody pressed Deploy.
- **Reports:** REP-001, REP-002

DEC-004 · PDF exports with pandoc and the existing Playwright browser

- **Question:** Which tool builds the decision PDFs?
- **Choice:** pandoc (Markdown to HTML) plus the project's existing Playwright Chromium (HTML to PDF). No installs.
- **Rationale:** Both are already available; no new dependency.
- **Player impact:** None.
- **Technical impact:** PDF export runs locally, not in CI.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Proposed
- **Delivery:** Implemented
- **Review trigger:** If pandoc is unavailable on the machine that exports.
- **Reports:** REP-001

DEC-005 · Budget interpretation: RM500 per month across all games

- **Question:** How should the RM500 monthly ceiling be applied?
- **Choice:** Treat it as portfolio-wide. Fruit Market's current spend is unknown (hypothesis: Cloudflare Free, or Workers Paid at about USD 5).
- **Rationale:** The workflow guide's conservative default.
- **Player impact:** None.
- **Technical impact:** Paid services need a quote and approval.
- **Data and security:** None.
- **Cost:** Unknown; to be recorded.
- **Owner:** Murshidi
- **Approval:** Proposed
- **Delivery:** Not started
- **Review trigger:** When the owner states the Cloudflare plan and other paid accounts.
- **Reports:** REP-001

DEC-006 · Preview = separate staging Worker

- **Question:** How should changes be previewed, given that Cloudflare preview and version URLs are not generated for Workers with Durable Objects?

- **Choice:** A separate Worker, fruit-market-staging, with its own rooms, leaderboard and stats, updated by CI after checks pass on main.
- **Rationale:** Closest free option to a per-change preview URL; keeps the live data untouched.
- **Player impact:** None; players never see the preview.
- **Technical impact:** A second Worker and a preview job in CI.
- **Data and security:** Separate storage; preview link is unlisted, not private.
- **Cost:** No new paid service; tiny usage within the existing plan.
- **Owner:** Murshidi
- **Approval:** Approved: “Confirm DEC-006, and plan PLAN-002” (2026-10-26)
- **Delivery:** Verified (PLAN-001)
- **Review trigger:** If Cloudflare adds preview URLs for Durable Object Workers.
- **Reports:** REP-002, REP-003

DEC-007 · Package B scope = PLAN-002

- **Question:** What should the records package contain?
- **Choice:** Canonical JSON records, generated views (register, changelog, state block, PDFs), a records check in verify, release tagging on Deploy, and migration of existing records.
- **Rationale:** Records were kept by hand in several places with nothing checking that they agreed.
- **Player impact:** None directly; releases become traceable.
- **Technical impact:** New records, a generator and checker script, and a deploy step.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-002 only, including its listed tests and documentation” (2026-10-26)
- **Delivery:** Verified (PLAN-002)
- **Review trigger:** If keeping records costs more than it saves.
- **Reports:** REP-003, REP-004

DEC-008 · Canonical record format: JSON records plus Markdown narrative reports

- **Question:** Where is the single source of truth for decisions and packages?
- **Choice:** JSON files in docs/workflow/records; hand-written REP reports stay as narrative snapshots; Markdown, changelog and PDF views are generated.
- **Rationale:** Node reads JSON without dependencies and the check can be strict. YAML needs a parser; plain Markdown can't be validated; a database is ruled out by the guide.
- **Player impact:** None.
- **Technical impact:** Claude edits the JSON; the owner reads the generated views.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-002 only, including its listed tests and documentation” (2026-10-26)
- **Delivery:** Verified (PLAN-002)
- **Review trigger:** If hand edits to JSON become a frequent source of errors.
- **Reports:** REP-003, REP-004

DEC-009 · Records check runs in verify and CI

- **Question:** Should record disagreements fail the build?
- **Choice:** Yes: npm run verify runs records:check, so CI fails on broken records.
- **Rationale:** Once branch protection is on, broken records can't be merged.
- **Player impact:** None.
- **Technical impact:** verify is about 1 second slower.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-002 only, including its listed tests and documentation” (2026-10-26)
- **Delivery:** Verified (PLAN-002)
- **Review trigger:** If the check blocks urgent fixes for record reasons.
- **Reports:** REP-003, REP-004

DEC-010 · Reconnect fix scope = PLAN-004 (evidence first, then fix)

- **Question:** How should the intermittent rejoin failure be fixed?
- **Choice:** Collect evidence and reproduce it with a stress test before changing game code; fix only the measured root cause, with a regression test.
- **Rationale:** Reading the code didn't identify a certain cause, and two failures in five runs are too rare to debug by guessing. The workflow guide says to reproduce and isolate before fixing.
- **Player impact:** Fewer stuck screens after a reload or dropped connection, once the fix is published.
- **Technical impact:** New stress test and diagnostics; a small fix in the online client or room server.
- **Data and security:** Diagnostics stay in CI artifacts for 7 days; they contain test players only.
- **Cost:** RM0 (GitHub Actions minutes are free for public repositories)
- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-004 only, including its listed tests and documentation” (2026-10-26)
- **Delivery:** Verified (PLAN-004)
- **Review trigger:** If not reproduced after 300 cycles: stop and report.
- **Reports:** REP-005, REP-006

DEC-011 · Stress test runs on demand, not on every push

- **Question:** Should the long reconnect stress test run in every CI run?
- **Choice:** No: a manual CI workflow plus local runs. Every-push CI keeps the existing tests (about 15 minutes).
- **Rationale:** Hundreds of cycles would make every CI run far slower; the stress test is for investigation and before releases.
- **Player impact:** None.
- **Technical impact:** One extra manual workflow.
- **Data and security:** None.
- **Cost:** RM0

- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-004 only, including its listed tests and documentation” (2026-10-26)
- **Delivery:** Verified (PLAN-004)
- **Review trigger:** If reconnect bugs slip through between releases.
- **Reports:** REP-005, REP-006

DEC-012 · 4-player clarity scope = PLAN-005 (screen text only)

- **Question:** How should the game make 4-player play discoverable?
- **Choice:** Clearer screen text and open-seat rows in the lobby, short hints on the home, quick-match and new-game screens, plus an automated 4-person test. No rule or server changes.
- **Rationale:** 4-player play already works; only discoverability is missing. Text changes are low risk and fast to verify.
- **Player impact:** Groups of 3 or 4 can find how to play together.
- **Technical impact:** src/ui/Online.tsx, src/ui/App.tsx, styles; one new browser test.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-005 only, including its listed tests and documentation” (2026-10-26)
- **Delivery:** Verified (PLAN-005)
- **Review trigger:** If playtesters still miss how to invite more than one friend.
- **Reports:** REP-007, REP-008

DEC-013 · Publish version 1.1.1 to players

- **Question:** Should the 1.1.1 candidate (reconnect fix, 4-player clarity, release process) replace 1.1.0 on the live game?
- **Choice:** Recommended after the owner's 15-minute preview playtest passes. Publication happens only when the owner runs Deploy LIVE game and types PUBLISH.
- **Rationale:** Fixes a bug players can hit (stuck on Connecting after a reload or rejoin) and makes group play discoverable. The game-code change is small (4 files, about 40 lines), the server code is unchanged, and all automated checks pass.
- **Player impact:** Fewer stuck online screens; clearer 2 to 4 player options.
- **Technical impact:** Live Worker updated to the candidate commit and tagged v1.1.1.
- **Data and security:** No data or storage format change; rolling back is safe.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “yes to 1 and 2 (1 = release 1.1.1: approve DEC-013, then the owner presses Deploy LIVE game)” (2026-10-26)
- **Delivery:** Not started
- **Review trigger:** If the preview playtest finds a problem, or the post-deploy smoke test fails.
- **Reports:** REP-009, REP-013

DEC-014 · Use the Concept and Theme Companion for creative work

- **Question:** Should Fruit Market's creative decisions follow the Game Concept and Theme Companion (edition 1.0) alongside the development workflow?
- **Choice:** Yes, with one shared record system (same reports, register and PDFs); creative approval never implies implementation approval.
- **Rationale:** Keeps creative and development decisions in one history and avoids a parallel register.
- **Player impact:** None directly.
- **Technical impact:** None; documentation process only.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Proposed
- **Delivery:** In progress
- **Review trigger:** If creative work starts needing its own package type.
- **Reports:** REP-010

DEC-015 · Adopt three design pillars

- **Question:** Which design pillars guide Fruit Market's creative choices?
- **Choice:** 1) Fresh fruit, fast money. 2) Outprice rivals in secret. 3) A cheerful town market. Each with inclusion and exclusion rules (REP-010 section 3).
- **Rationale:** Derived from the actual rules and the strongest theme-to-mechanic matches found in the review.
- **Player impact:** Future features and content are judged against these.
- **Technical impact:** None now.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “yes to 1 and 2 (2 = approve DEC-015 and DEC-016)” (2026-10-26)
- **Delivery:** Implemented
- **Review trigger:** If playtests show players value something the pillars exclude.
- **Reports:** REP-010, REP-011, REP-013

DEC-016 · First creative test: reveal-comprehension playtest

- **Question:** What is the smallest useful test of the riskiest creative hypothesis?
- **Choice:** 3 to 5 first-time players, Quick game vs 1 computer shop, no coaching; continue if at least 3 of 4 can explain a lost sale by day 3 (provisional, owner-set). No code.
- **Rationale:** Pillar 2 depends on players understanding why they won or lost a sale; no player evidence exists yet.
- **Player impact:** Testers play the existing game.
- **Technical impact:** None.
- **Data and security:** Notes only; no personal data kept.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “yes to 1 and 2 (2 = approve DEC-015 and DEC-016)” (2026-10-26)
- **Delivery:** Not started
- **Review trigger:** After the first playtest round.

- **Reports:** REP-010, REP-011, REP-013

DEC-017 · Setting: a market where the whole world shops

- **Question:** Where is Fruit Market's market?
- **Choice:** Option C: a deliberate 'the whole world shops here' mix. A sunny harbour market; customers from every culture and fantasy role; seaside palette; plain market dollars.
- **Rationale:** Keeps the existing art and cast, turns the mix into an intentional identity, and costs nothing now.
- **Player impact:** No visible change now; guides future content.
- **Technical impact:** None now; recorded in the creative bible.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “1. Where is this market? C: a deliberate “the whole world shops here” mix.” (2026-10-26)
- **Delivery:** Implemented
- **Review trigger:** If playtesters find the setting confusing.
- **Reports:** REP-011

DEC-018 · Rename Panda Lim while keeping the cast coherent

- **Question:** Should the 'Panda Lim' rival be renamed?
- **Choice:** Yes: an animal nickname that fits beside Foxy Rosa, Big Benny, Uncle Toad, Rocky and Pinky; picture and lines unchanged. Cast rules recorded in the creative bible.
- **Rationale:** Avoids pairing a surname with a national symbol while preserving the animal-trader theme.
- **Player impact:** New games show the new name.
- **Technical impact:** One data line plus tests (PLAN-006).
- **Data and security:** Saved games keep the old name until they end.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “Rename Panda Lim? - sure, but dont run away from the general theme coherence” (2026-10-26)
- **Delivery:** Verified (PLAN-006)
- **Review trigger:** If the new name tests poorly with players.
- **Reports:** REP-011, REP-012

DEC-019 · Rename scope = PLAN-006, shipped in 1.1.2

- **Question:** How and when should the rename be built?
- **Choice:** As PLAN-006 (display name, reserved names, compatibility tests), released as 1.1.2 after 1.1.1 is live.
- **Rationale:** Keeps the 1.1.1 candidate unchanged; the rename is small and safe but still needs its own checks.
- **Player impact:** New name from 1.1.2.
- **Technical impact:** Small data change plus tests.
- **Data and security:** None.

- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Approved: “Approve PLAN-006 only, including its listed tests and documentation, name: Patches” (2026-10-26)
- **Delivery:** Verified (PLAN-006)
- **Review trigger:** If 1.1.1 is delayed for long.
- **Reports:** REP-011, REP-012

DEC-020 · UI and UX polish in three phases

- **Question:** How should the game's look and feel be made to feel finished?
- **Choice:** Phase 1 fixes 9 visible defects (PLAN-007, ships in 1.1.2); phase 2 calms the table layout and phase 3 adds a design system, motion and sound, both designed as phone-viewable mockups first (PLAN-008) and built only after the owner picks; playtest before and after.
- **Rationale:** The review found concrete causes (defects, clutter, no focus, inconsistent system). Fixing defects is cheap and safe; layout changes need the owner's eye and player evidence.
- **Player impact:** A cleaner, clearer game in steps.
- **Technical impact:** Presentation code only; no rules, saves or server.
- **Data and security:** None.
- **Cost:** RM0
- **Owner:** Murshidi
- **Approval:** Awaiting approval
- **Delivery:** Not started
- **Review trigger:** If the baseline playtest shows a different priority.
- **Reports:** REP-014, REP-015

DEC-021 · Art direction for the polish: explore A/B/C, recommend toy-box tabletop

- **Question:** Which art direction should the polish follow, given the owner dislikes the current art style and points to It Takes Two?
- **Choice:** Run a style test of A (polished pixel), B (toy-box tabletop 2D) and C (3D toy diorama) on one screen; recommendation B, drawn by Claude as vector art, with a paid illustrator as a later, separately priced option.
- **Rationale:** B addresses clutter, feel and art style, keeps phones fast and costs nothing to test; C is too large and risky now; A does not answer the art-style concern.
- **Player impact:** A warmer, toy-like look with more motion, if chosen and built.
- **Technical impact:** Redraw of fruits, about 30 characters, rivals, about 40 icons and the UI (only after approval of a build plan).
- **Data and security:** None.
- **Cost:** RM0 for the test; any illustrator needs a quote and approval
- **Owner:** Murshidi
- **Approval:** Awaiting approval
- **Delivery:** Not started (PLAN-008)
- **Review trigger:** After the style test, or if the owner wants a paid illustrator.
- **Reports:** REP-015

History (oldest first)

- 2026-10-26 13:15 · REP-001: baseline audit; PLAN-001 proposed; DEC-001 to DEC-005 proposed.
- 2026-10-26 13:20 · Owner: "Approve PLAN-001 only, including its listed tests and documentation". DEC-002 and DEC-003 approved; PLAN-001 approved.
- 2026-10-26 13:25 · PLAN-001 implemented on main (2e4b240). Push started CI and Android only; no live deploy.
- 2026-10-26 13:42 · REP-002: PLAN-001 verified except A4 (owner playtest). DEC-006 recorded, awaiting confirmation.
- 2026-10-26 13:46 · Owner: "Confirm DEC-006, and plan PLAN-002". DEC-006 approved. REP-003: PLAN-002 proposed; DEC-007 to DEC-009 awaiting approval.
- 2026-10-26 13:48 · Owner: "Approve PLAN-002 only, including its listed tests and documentation". PLAN-002, DEC-007, DEC-008 and DEC-009 approved.
- 2026-10-26 13:50 · Correction: the report headers of REP-002 (13:45) and REP-003 (13:50) were rounded estimates; the commit times 13:42 and 13:46 are used above. The reports are left as written.
- 2026-10-26 13:52 · PLAN-002 implemented on workflow/plan-002: records migrated; B1, B2, B3, B5 and B7 PASS; B4 and B6 pending.
- 2026-10-26 14:09 · PLAN-002 verified: B1 to B7 PASS (CI run 40276699547 on a7fa577). DEC-007, DEC-008 and DEC-009 delivered.
- 2026-10-26 14:09 · Known issue found: CI run 40275680793 (docs-only commit 63c3682) failed because the online test's page reload did not show the game within 20 s. Same symptom as the quick-match failure in CI run 1. Pre-existing and intermittent (2 of 5 CI runs); not part of PLAN-002. Proposed as the next package in REP-004.
- 2026-10-26 14:16 · Owner: "PLAN the reconnect fix". REP-005: PLAN-004 proposed; DEC-010 and DEC-011 awaiting approval.
- 2026-10-26 14:18 · Owner: "Approve PLAN-004 only, including its listed tests and documentation". PLAN-004, DEC-010 and DEC-011 approved. (Time approximate: between the REP-005 commit at 14:17 and the first work commit at 14:27.)
- 2026-10-26 14:53 · PLAN-004: root cause reproduced (lost first update after rejoining; 20 of 20 hidden-tab reloads stuck before the fix) and fixed in the client (9d9e911).
- 2026-10-26 15:01 · PLAN-004: R1 to R5 PASS (165 hidden-tab cycles after the fix, 0 failed). R6 waiting for 3 CI runs.
- 2026-10-26 15:34 · REP-006: PLAN-004 verified (R1 to R6 PASS). The fix is on the preview; the live game is unchanged until a release is prepared and deployed.
- 2026-10-26 18:34 · Owner asked how to play with 4 people and requested a plan. Check: 4-person online room works locally (4 joined and played; a 5th refused). REP-007: PLAN-005 proposed; DEC-012 awaiting approval.
- 2026-10-26 18:36 · Owner: "Approve PLAN-005 only, including its listed tests and documentation". PLAN-005 and DEC-012 approved.
- 2026-10-26 18:44 · PLAN-005 implemented: C1 to C5 PASS; C6 waiting for CI.
- 2026-10-26 19:01 · REP-008: PLAN-005 verified (C1 to C6 PASS). On the preview; goes live with the next release.
- 2026-10-26 19:02 · Correction: the REP-008 header time (19:05) was an estimate; the report was completed at 19:01. The report is left as written.

- 2026-10-26 20:26 · Owner: "PREPARE RELEASE". Version set to 1.1.1 (package and Android versionName); release candidate entry added; DEC-013 (publish 1.1.1) awaiting approval. REP-009.
- 2026-10-26 20:27 · CI run 13 (f472186, records-only commit) was cancelled: the browser download took 16 minutes, so the job hit its 20-minute limit. Not a code failure. Limit raised to 35 minutes in the 1.1.1 release commit.
- 2026-10-26 21:47 · Owner asked to apply the Game Concept and Theme Companion. REP-010: theme review (no changes made); DEC-014 proposed; DEC-015 (pillars) and DEC-016 (first playtest) awaiting approval. Open questions: setting identity, renaming Panda Lim, availability of 3 to 5 testers.
- 2026-10-26 22:14 · Owner answers to REP-010: setting C approved (DEC-017); rename Panda Lim approved, keeping the cast coherent (DEC-018); 3 to 5 first-time testers are available. REP-011: creative bible written; PLAN-006 (rename) and DEC-019 awaiting approval. DEC-015 and DEC-016 still await explicit approval.
- 2026-10-26 22:17 · Owner: "Approve PLAN-006 only, including its listed tests and documentation, name: Patches". PLAN-006 and DEC-019 approved; new name Patches.
- 2026-10-26 22:18 · PLAN-006 implemented on workflow/plan-006: N1 to N4 PASS (5 new tests); N5 waiting for CI.
- 2026-10-26 22:48 · REP-012: PLAN-006 verified (N1 to N5 PASS) on workflow/plan-006. Waits off main until 1.1.1 is live; then ships as 1.1.2.
- 2026-10-26 22:56 · Owner: "yes to 1 and 2". DEC-013 (publish 1.1.1) approved; the owner runs Deploy. The owner's own preview playtest result was not reported (recorded as NOT RUN). DEC-015 (pillars) and DEC-016 (first playtest) approved. Recorded on workflow/plan-006 so main stays on the 1.1.1 candidate.
- 2026-10-26 22:57 · Correction: the REP-013 header time (22:58) was an estimate; the report was written at 22:56. The report is left as written.
- 2026-10-26 23:06 · Owner: the game doesn't feel polished; propose changes on the whole UI and UX. REP-014: review of all main screens at 360×640 and 390×844; DEC-020, PLAN-007 (phase 1 fixes) and PLAN-008 (mockups) awaiting approval. PDF export now inlines report images.
- 2026-10-26 23:18 · Owner answers to REP-014: most bothered by clutter, lack of animation and feel, and the art style itself; polished reference: It Takes Two. REP-015: art-direction options; PLAN-008 revised to a style test plus mockups (still awaiting approval); DEC-021 awaiting approval; PLAN-007 suggested to shrink to 3 fixes. PLAN-007 and PLAN-008 not yet approved.

REP-001 · Baseline audit and PLAN-001 (Package A)

Game	Fruit Market (`murshidihashim/Fruit-Market`)
Report date	2026-10-26 13:15 (Asia/Kuala_Lumpur)
Build audited	`main @ c915e49` (CI run 2 green, deployed)
Prepared by	Claude (studio lead), from the "Personal workflow for Murshidi", edition 1.2
Type	Planning only. **No game implementation changes were made.**
Approval status	PLAN-001 **awaiting approval**

1 · Baseline: where the game stands

Stage: between *MVP / private alpha* and *browser beta*. The game is complete and publicly reachable, but no unfamiliar players have been observed yet, so player understanding and fun are unproven.

Working and must be preserved

- Rules engine: deterministic, seeded, a single entry point (`applyAction`), rules version 2. 63 unit tests, including 200 simulated games.
- Saves: versioned envelope (`v2` plus rules version); old saves still load.
- Online play is server-authoritative (Cloudflare Worker with Durable Objects):
  - secret prices are redacted; the leak check reads every network message and found 0 leaks
  - reconnect, a 75 s grace period then computer takeover, and quick match
  - per-IP rate limits, a name policy, and a replay-verified daily leaderboard
- Android build (Capacitor) is produced by CI on every push.
- Docs: `RULES.md` , `ONLINE.md` , `PUBLISHING.md` , `docs/ANALYTICS.md` , the studio audit `docs/studio/01–10` (backlog = `10`).

Gaps against the workflow guide

#	Finding	Type	Severity
G1	Every push to `main` deploys straight to the public URL. The guide requires a separate publication approval.	Process defect	High
G2	`main` has **no branch protection** (checked via GitHub: `protected: false`). CI results don't block anything.	Measured	High
G3	No single verification command. There are separate `lint` , `typecheck` , `test` , `build` , `test:e2e` scripts.	Measured	Medium
G4	No `CLAUDE.md` , current-state record, game brief, or owner brief. Each session relies on long chat history.	Measured	Medium
G5	No decision records or PDF register before this report.	Measured	Low
G6	Real-device testing (Android Chrome, iPhone Safari): never done.	NOT RUN	High for release

#	Finding	Type	Severity
G7	No backup or restore for server data (leaderboards, stats). Room crash recovery not tested.	NOT RUN	Medium
G8	Monthly hosting cost unknown (Cloudflare plan not recorded).	Unknown	Medium
G9	The game-screen open time is noisy (217–489 ms locally). Quick match failed once in CI and passed on the rerun.	Measured, pre-existing	Low

2 · Readiness baseline (evidence as of today)

Gate	Check	Result	Evidence and limitation
Build ready	Clean install, lint, typecheck, unit tests, build	PASS	CI run 40266399588 on `c915e49`
Play ready	Automated solo game to results	PASS	`scripts/smoke.cjs` , local and CI
Play ready	Owner completes a match in a preview	NOT RUN	Needs you
Online ready	Two independent clients finish a match; results agree	PASS	`e2e-online.cjs` , local and CI run 2
Online ready	Private state: no secret prices in any message	PASS	40 frames checked, 0 leaks
Online ready	Illegal actions, wrong phase, votes, takeover	PASS	`tests/room.test.ts` (16 tests)
Online ready	Reconnect, abandonment, quick match	PASS	e2e; one CI flake, then a pass
Online ready	Server crash recovery	NOT RUN	Room state is persisted per action; never tested
Release ready	Real Android Chrome and iPhone Safari	NOT RUN	Emulation only
Release ready	Backup and restore, rollback drill	NOT RUN	—
Release ready	Live smoke test after deploy	NOT RUN	The sandbox proxy blocks the live URL; the deploy job itself succeeds
Release ready	Publication approval gate	FAIL	See G1

Release-ready is **not claimed**.

3 · PLAN-001 · Package A: baseline, one verification path, preview and owner brief

Player problem: updates reach players without your approval and without one clear pass/fail check, so a bad change could ship silently. You also have no short guide for trying a build yourself.

Included scope

1. **One verification command.** `npm run verify` runs lint, typecheck, unit tests and build, and stops at the first failure (about 1 minute). `npm run verify:full` adds the browser suite (about 15 minutes). CI calls `verify:full`. Existing scripts are reused; nothing new is installed.
2. **Preview before publication.**
   - Pushes to `main` run checks only.
   - Production deploys happen only when you trigger *Deploy* (a manual GitHub button), which first runs `verify`.
   - Each change also gets a private Cloudflare **preview URL** that doesn't touch the live game. I'll confirm in Cloudflare's current docs that this is free on your plan; if it isn't, the local preview steps are the fallback.
3. **Branch protection instructions.** A 5-step guide for you to turn it on in GitHub settings (I can't change repository settings). It requires the CI check before merging.
4. **Short records** (reusing the studio docs, not duplicating them):
   - `CLAUDE.md`, about 90 lines: rules, actual commands, preservation constraints, pointers
   - `docs/workflow/STATE.md`, about 500 words
   - `docs/workflow/GAME-BRIEF.md`: audience hypothesis, core loop, session length, exclusions
5. **Owner brief.** `docs/workflow/OWNER-BRIEF.md`: how to open a preview, a 10-minute playtest checklist, and what each check means in plain English.

Excluded: any change to rules, balance, art, saves, the server protocol or hosting; paid services; publication; Package B (record automation) and Package C (devices, backups, capacity).

Roles: producer (scope), DevOps and release (deploy gate, CI), QA (verify command), technical lead (`CLAUDE.md`). Self-review only; low risk.

Likely files: `package.json` (scripts), `.github/workflows/{ci,deploy-web}.yml`, new `CLAUDE.md`, new `docs/workflow/{STATE,GAME-BRIEF,OWNER-BRIEF}.md`, `README.md` (pointer).

Acceptance (observable)

- A1. `npm run verify` exits 0 on the current code. A deliberately broken test makes it exit non-zero (shown, then reverted).
- A2. Pushing to `main` does **not** change the live site. Running *Deploy* does, and only after `verify` passes.
- A3. A preview URL (or documented local preview) lets you play the build before you approve publication.
- A4. You can follow `OWNER-BRIEF.md` to open a preview and finish a match, without asking me anything.
- A5. `CLAUDE.md` stays under 120 lines; `STATE.md` stays at or under about 600 words.

Checks: `verify` (PASS and FAIL demonstration), one CI run, one manual deploy dry run (after your separate publication approval, or skipped and marked NOT RUN), and the line and word counts.

Cost and data: RM0 extra. No player data is touched.

Risks: a preview URL feature may need a Cloudflare setting (fallback: local preview). If you forget to press *Deploy*, fixes wait (that's the point).

Rollback: revert the package commit; the deploy workflow returns to deploying on every push.

4 · Packages B and C (outline only; detailed later)

- **PLAN-002 · Package B, records and PDFs:**
  - one canonical record per package (Markdown with an ID header)
  - a deterministic local script that builds the owner brief, changelog, report PDFs and cumulative register, using pandoc and the project's existing Playwright browser (no new dependency)
  - export failures are labelled PDF EXPORT BLOCKED
- **PLAN-003 · Package C, online, capacity and release evidence:**
  - real-device checklist runs (you plus your phones)
  - Durable Object data export and restore drill
  - a room crash-recovery test
  - a small load test within an approved budget, with written workload assumptions (peak concurrent players, 2–4 per room, about 15 min matches)
  - cost record with a dated Cloudflare quote
  - live smoke test from an unrestricted machine (CI)

5 · Decisions

ID	Question	Recommendation	Status	Delivery
DEC-001	Adopt the workflow guide 1.2 for Fruit Market?	Yes; one package at a time	Proposed	Not started
DEC-002	What goes in Package A?	PLAN-001 scope above	Awaiting approval	Not started
DEC-003	When does production deploy?	Only on your manual *Deploy* after `verify`; never on push	Awaiting approval	Not started
DEC-004	PDF tooling	pandoc plus the existing Playwright browser; no installs	Proposed	Used for this report
DEC-005	Budget interpretation	RM500/month portfolio-wide ceiling; current Fruit Market spend unknown (hypothesis: Cloudflare Free, or Workers Paid at about USD 5)	Proposed (assumption)	Not started

Rationale for DEC-003: it is the only way to make "a preview is not a release" true for this repo. Trade-off: one extra tap per release.

6 · Questions for you (max 3)

1. **Devices:** which real phones can you test on (Android model, iPhone model)? This decides Package C.
2. **Hosting cost:** is your Cloudflare account on the Free plan or Workers Paid? Any other paid accounts yet (domain, Google Play developer account)?
3. **Approve PLAN-001** as written (including DEC-003)? Or change it?

7 · Sources

- Repository inspection, `main @ c915e49`, 2026-10-26.
- GitHub branch API (`main protected: false`), 2026-10-26.
- CI run 40266399588 (green), deploy run 40266399518 (success), Android run 40266399555 (success).
- Workflow guide: "Personal workflow for Murshidi", consolidated edition 1.2, 26 Oct 2026.
- Preview-URL and plan-pricing facts to be confirmed against current Cloudflare docs during PLAN-001 (not yet verified).

REP-002 · PLAN-001 (Package A) implemented and verified

Game	Fruit Market
Report date	2026-10-26 13:45 (Asia/Kuala_Lumpur)
Package	PLAN-001, approved by the owner 2026-10-26: "Approve PLAN-001 only, including its listed tests and documentation"
Build	`main @ 2e4b240` ; records on `workflow/plan-001`
Live game	**Unchanged** (still `c915e49`). No publication approval was given and none was performed.
Delivery status	Implemented; verified except A4, which needs the owner

What changed

Added

- `npm run verify` : lint, types, unit tests, build; stops at the first failure.
- `npm run verify:full` : `verify` plus browser tests.
- A **preview** copy of the game, updated automatically after CI passes on `main` , then smoke-tested: https://fruit-market-staging.murshidi-hashim.workers.dev. It is a separate Worker with its own rooms, leaderboard and stats.
- `scripts/smoke-live.mjs` : after a deploy, checks that the page loads and the server runs the expected rules version.
- `CLAUDE.md` (72 lines), `docs/workflow/STATE.md` , `GAME-BRIEF.md` , `OWNER-BRIEF.md` (preview, playtest checklist, publishing, branch-protection steps).

Changed

- The live deploy now runs only through GitHub Actions → *Deploy LIVE game* → type `PUBLISH` . It runs `verify` first and a live smoke test after.
- CI runs `verify:full` . README and ONLINE.md describe the new path.
- Lint now covers the docs export script (it found 12 errors in it, which `verify` caught; now fixed).

Removed: automatic live deploys on every push to `main` .

Player impact: none yet. Players only benefit once you publish, because broken changes can no longer reach them silently.

Acceptance results

#	Criterion	Result	Evidence
A1	`verify` exits 0 on current code; non-zero with a broken	PASS	Local: exit 0 (63 tests, build), then exit 1 with one test deliberately broken ("1 failed, 62 passed"), then exit 0

#	Criterion	Result	Evidence
	test		once restored
A2	Push to `main` does not change the live site	PASS	Push of `2e4b240` started only *CI* (run 40274184197) and *Android build*; no live deploy run
A2	*Deploy* runs only after `verify` passes	PASS (by design) / NOT RUN (live)	`verify` is the step before deploy; not run because publication isn't approved
A3	A preview lets you play before publishing	PASS	Preview job: "Deployed fruit-market-staging", then "SMOKE OK … rules version 2" (version `adf6c304`)
A4	Owner can open the preview and finish a match using the brief	NOT RUN	Needs you (OWNER-BRIEF §1–2)
A5	`CLAUDE.md` under 120 lines; STATE about 500–600 words	PASS	72 lines; STATE about 350 words
—	Full CI (`verify:full`) on `2e4b240`	PASS	Checks job green in 14 min: lint, types, 63 tests, build, solo game, performance budget, online match with a leaver (0 leaks), quick match
—	Smoke script fails on a dead address	PASS	Local: "SMOKE FAILED … fetch failed", exit 1
—	Preview online rooms (Durable Objects) in a real browser	NOT RUN	The smoke test checks the page and API only. Covered by the optional online step in the playtest checklist

Deviation from the plan (needs your confirmation)

PLAN-001 said "a private Cloudflare preview URL; if it isn't free, local preview steps". Cloudflare's docs state that preview and version URLs **are not generated for Workers that implement a Durable Object**, and Fruit Market's server uses Durable Objects for rooms. So I used the next closest option: a second, separate Worker (`fruit-market-staging`). This is recorded as **DEC-006**.

- Cost: no new paid service. It runs within the Cloudflare plan you already use; usage is tiny. Your plan type is still unknown (DEC-005).
- "Private" means unlisted: anyone with the link can open it.
- To undo: delete the `fruit-market-staging` Worker in Cloudflare and remove the `preview` job.

Risks and costs

- Expected extra cost: RM0 (no new paid service). Not yet confirmed against your plan type.
- Until branch protection is on, a direct push can still skip CI, but it can no longer reach players without *Deploy*.
- If CI is red, the preview isn't updated; the live game is unaffected.

Decisions

ID	Decision	Approval status	Delivery
DEC-002	Package A scope = PLAN-001	Approved (owner, 2026-10-26)	Verified (A4 pending)
DEC-003	Live deploys only via manual *Deploy* after `verify`	Approved (owner, 2026-10-26, as part of PLAN-001)	Verified
DEC-006	Preview = separate staging Worker, because per-version URLs don't support Durable Objects	Awaiting confirmation	Implemented

Sources

- Cloudflare docs, Version URLs and Preview URLs (`developers.cloudflare.com/workers/versions-and-deployments/version-urls/` and `/preview-urls/`): not generated for Workers that implement a Durable Object. Retrieved 2026-10-26 via search.
- Wrangler config reference vs the Durable Objects environments page disagree on per-environment migrations; resolved empirically: the staging deploy with top-level migrations succeeded.
- CI run 40274184197 (checks and preview jobs green).

REP-003 · DEC-006 confirmed; PLAN-002 (Package B) proposed

Game	Fruit Market
Report date	2026-10-26 13:50 (Asia/Kuala_Lumpur)
Build	`main @ 63c3682` (live game still `c915e49`)
Type	Planning only. **No game implementation changes were made.**
Approval status	DEC-006 **approved** (owner, 2026-10-26: "Confirm DEC-006"). PLAN-002 **awaiting approval**

1 · DEC-006 confirmed

The preview stays a separate staging Worker (`fruit-market-staging`). Approval status: approved; delivery: verified (REP-002, A3).

2 · Problem Package B solves

Records exist, but they're kept by hand:

- Each decision's status is typed in two or three places (the register, every report that mentions it, STATE). Nothing checks that they agree. REP-002 already needed a manual word-count correction.
- Nothing stops an approval being recorded without your words as evidence, or history being edited.
- There's no player-facing changelog and no version tags, so "what changed in this release" can't be answered from the repo.
- The PDF export works (built in REP-001) but isn't part of any check.

Players never see this directly. It matters because it keeps approvals honest and makes releases traceable when there's money or player data involved later.

3 · PLAN-002 · Package B: canonical records, generated views, checks

Included scope

1. **One canonical source per record** (JSON, no new tools):
   - `docs/workflow/records/decisions.json`: every DEC with question, choice, rationale, impacts, owner, approval evidence (your exact words and date), delivery status, links, review trigger, supersession. A `history` list is only ever appended to.
   - `docs/workflow/records/PLAN-###.json`: one per package with scope, exclusions, acceptance, approval evidence, changes (Added / Changed / Fixed / Removed), checks, build, release state.
   - Reports (`reports/REP-###.md`) stay as hand-written narrative snapshots and are never rewritten.
2. **Generated views**, built by `npm run records` with a deterministic script:

    - `DECISIONS.md`: current decisions and pending approvals first, then the full history.
    - `CHANGELOG.md`: player-facing, grouped by release, with "unreleased" vs "published" state. It is seeded with one summary entry for everything before the workflow, taken from git history.
    - The "Pending approvals" block inside `STATE.md` (between markers; the rest stays hand-written).
    - PDFs: each report, plus the cumulative register (current decisions, then pending approvals, then all reports in order).
3. **Checks**, via `npm run records:check`, added to `npm run verify` so CI runs it. It fails when:
    - an ID is duplicated or skipped, or a report mentions a DEC, PLAN or REP that doesn't exist
    - a decision is "approved" without owner evidence, or a delivery status isn't one of the allowed values
    - a history entry was changed or removed compared with the last commit (append-only)
    - a generated view is out of date with its source
4. **Release tagging**: when you publish, the deploy records the release in the changelog and tags the commit (`v1.1.1` and so on), so every live build can be traced. This only runs when you press *Deploy*.
5. **Migration**: move DEC-001 to DEC-006, PLAN-001 and REP-001 to REP-003 into the new records without losing anything.

Excluded

- Game code, rules, server and hosting.
- Package C (devices, backups, capacity).
- A cross-game studio register (NOT APPLICABLE until there's a second game).
- Installing anything. PDF building stays local because it uses pandoc. If pandoc is missing, the script prints PDF EXPORT BLOCKED and keeps the last good PDFs; it isn't part of CI.

Roles: producer (record design), technical lead (scripts), QA (failure cases), DevOps (tagging). Self-review only; low risk.

Likely files: new `docs/workflow/records/*.json`, new `scripts/records.mjs`, new `CHANGELOG.md`, `docs/workflow/export-pdf.cjs` (reads the records), `package.json` (scripts), `.github/workflows/deploy-web.yml` (tag and changelog step), `CLAUDE.md` and `OWNER-BRIEF.md` (one line each).

Acceptance (observable)

- B1. `npm run records` rebuilds all views. Running it twice in a row changes nothing (no git diff in the Markdown outputs).
- B2. `records:check` fails, with a plain-English message, for each of these deliberate breaks (shown, then reverted):
    - a duplicate ID
    - an unknown DEC referenced in a report
    - "approved" without evidence
    - an edited history line
    - a stale `DECISIONS.md`

- B3. With pandoc made unavailable, the export prints PDF EXPORT BLOCKED, exits non-zero and leaves the previous PDFs untouched.
- B4. The cumulative PDF opens with current decisions and pending approvals, followed by REP-001 to REP-003 in order.
- B5. After migration, every decision, status and approval quote from REP-001 to REP-003 appears in the records.
- B6. `npm run verify` (and CI) runs `records:check` and stays green on the migrated records.
- B7. Tagging is tested on a throwaway branch: a dry run shows the tag and changelog entry it would create. The real tag happens only at your next publication.

Checks: B1–B7, plus `verify`, plus one CI run.

Cost and data: RM0. No player data. `npm run verify` gets about 1 second slower.

Risks:

- Record-keeping overhead: JSON is less friendly to edit by hand. Mitigation: Claude edits the records; you only read the generated Markdown and PDFs.
- The append-only check compares with the last commit, so a deliberate correction needs a new "correction" history entry rather than an edit. That is intended.

Rollback: revert the package commits. The hand-written `DECISIONS.md` and reports remain valid.

4 · Decisions

ID	Question	Recommendation	Status	Delivery
DEC-006	Preview = separate staging Worker	Keep	Approved (owner, 2026-10-26)	Verified
DEC-007	Package B scope = PLAN-002	As above	Awaiting approval	Not started
DEC-008	Canonical record format	JSON records plus Markdown narrative reports; Markdown, changelog and PDF views generated	Awaiting approval (part of PLAN-002)	Not started
DEC-009	Records check in `verify` and CI	Yes, so disagreements block merges once branch protection is on	Awaiting approval (part of PLAN-002)	Not started

Options considered for DEC-008:

- Keep hand-written Markdown only: free, but can't be validated reliably.
- YAML front matter: friendlier, but needs a parser dependency.
- A database: ruled out by the workflow guide.
- **JSON (recommended):** Node reads it without dependencies, and the check can be strict.

5 · Still open from earlier (not blocking PLAN-002)

- Which real phones you can test on (needed for Package C).
- Your Cloudflare plan (Free or Workers Paid) and any other paid accounts (DEC-005).
- Preview playtest (REP-002, A4) and branch protection (owner brief, section 5).

REP-004 · PLAN-002 (Package B) implemented and verified

Game	Fruit Market
Report date	2026-10-26 14:10 (Asia/Kuala_Lumpur)
Package	PLAN-002, approved by the owner 2026-10-26: "Approve PLAN-002 only, including its listed tests and documentation"
Build	`main @ a7fa577` (records), plus this report
Live game	**Unchanged** (still `c915e49`, version 1.1.0). No publication.
Delivery status	Verified (B1–B7 PASS)

What changed

Added

- `docs/workflow/records/`, the single source for every record:
  - `decisions.json`: all decisions, plus a history that can only grow
  - `PLAN-001.json` to `PLAN-003.json`
  - `releases.json`
- `npm run records`: rebuilds `DECISIONS.md`, `CHANGELOG.md`, the pending-approvals block in `STATE.md`, and every PDF.
- `npm run records:check` (now the first step of `verify`, so CI runs it). It rejects:
  - broken or duplicate IDs
  - references to things that don't exist
  - an approval without your words and a date
  - an edited history line or a rewritten approval
  - out-of-date views
- `CHANGELOG.md`, in plain player language. It is seeded with version 1.1.0 (what's live now).
- Release step in *Deploy LIVE game*: it refuses to run without release notes for the new version, and after a successful deploy it tags the commit (for example `v1.1.1`).

Changed

- The PDF export is all-or-nothing. On failure it prints PDF EXPORT BLOCKED and keeps the previous PDFs.
- CI fetches the full history so the records check can compare with the previous commit.

Removed

- The hand-edited `DECISIONS.md` (it is now generated).

Player impact: none directly. The release process is now traceable: every live build will have a version tag and changelog notes.

Acceptance results

#	Criterion	Result	Evidence
B1	Rebuilding twice changes nothing	PASS	Checksums of all three generated files identical after a second run
B2	Check fails on 5 deliberate breaks	PASS	Duplicate DEC-008; REP-003 citing a nonexistent DEC; DEC-001 "approved" with no evidence; history line 1 edited; stale register. Each exits 1 with a plain message. All reverted
B3	No pandoc → PDF EXPORT BLOCKED, old PDFs kept	PASS	Exit 1; all 4 PDFs byte-identical; no temporary files left
B4	Cumulative PDF order	PASS	Pending approvals, then current decisions, then REP-001, REP-002, REP-003
B5	Migration lost nothing	PASS	Every row of the old register matches. DEC-004 "In use" maps to the allowed value "implemented"; DEC-007 to DEC-009 changed only because of today's approval (in history)
B6	`verify` and CI run the check, green	PASS	CI run 40276699547: "RECORDS OK … history checked against 63c3682", full suite green, preview updated
B7	Tagging dry run	PASS	Throwaway branch: refuses while 1.1.0 is published; with a 1.1.1 candidate, "would tag v1.1.1" plus the notes. Branch deleted, no tag created

Known issue found (not part of this package)

The game screen sometimes doesn't appear after rejoining an online room. The automated tests caught it in 2 of the last 5 CI runs, with no game-code change between them:

- Run 1: after quick match, the second player's game didn't show within 15 s.
- Run 4 (a docs-only commit): after reloading the page mid-match, the game didn't come back within 20 s.

Every other run passed, including locally (more than 8 runs). Real players could hit this when their phone reloads the page or their connection drops.

This is pre-existing and intermittent; the cause isn't known yet. I recommend it as the next package, before Package C: reproduce it by repeating the reconnect test many times, find the cause, fix it, and prove the fix with repeated runs.

Risks and costs

- RM0. `verify` is about 1 second slower.
- The first live release under the new process needs "PREPARE RELEASE" first, because Deploy refuses without release notes. That is intended.
- Until the reconnect issue is fixed, CI can occasionally fail for reasons unrelated to the change being tested. Read red runs before assuming the latest change caused them.

Decisions

ID	Decision	Approval	Delivery
DEC-007	Package B scope = PLAN-002	Approved (owner, 2026-10-26)	Verified
DEC-008	JSON records plus narrative reports; generated views	Approved (owner, 2026-10-26)	Verified
DEC-009	Records check in `verify` and CI	Approved (owner, 2026-10-26)	Verified

Still pending your decision: DEC-001 (adopt the workflow), DEC-004 (PDF tooling), DEC-005 (budget assumption) and PLAN-003 (outline). See the register's "Pending approvals".

REP-005 · PLAN-004 proposed: fix the rejoin problem in online play

Game	Fruit Market
Report date	2026-10-26 14:20 (Asia/Kuala_Lumpur)
Build	`main @ 935a287` (live game still 1.1.0, `c915e49`)
Type	Planning only. **No game implementation changes were made.** Read-only code inspection only.
Approval status	PLAN-004 **awaiting approval** (DEC-010, DEC-011)

1 · The problem

Sometimes, after a player rejoins an online room, the game screen doesn't come back:

CI run	Situation	Waited	Result
Run 1 (`da2eea9`)	Second player joins via Quick match	15 s	Game screen never appeared
Run 4 (`63c3682`, docs only)	Player A reloads the page mid-match	20 s	Game screen never appeared
Runs 2, 3 and 5, plus 8+ local runs	Same steps	—	Passed (usually about 1 s)

There was no game-code change between the failing and passing runs, so the problem is timing-dependent. Real players can trigger it: phones reload pages when memory is low, and mobile connections drop. The tests saved **no evidence** when it happened: the local server's log was discarded, and the online test takes no screenshot on failure.

2 · What the code shows (read-only)

How rejoining works today:

1. The phone keeps a seat token per room.
2. On reload, the online screen sees `?room=CODE` plus the token and reconnects.
3. The server's `hello` matches the token to the seat and sends the room state.
4. The game screen appears once the state says the match has started.

Hypotheses, none confirmed yet:

#	Hypothesis	Where	Fits run 1	Fits run 4	Evidence so far
H1	In a quick-match lobby, a seat is deleted when its connection closes. If the waiting player's connection	`server/room.ts closed()`	Yes	No	Plausible from code

#	Hypothesis	Where	Fits run 1	Fits run 4	Evidence so far
	blips just as the stranger joins, the match never starts				
H2	The old connection's "closed" event reaches the server *after* the new `hello`, wrongly marking the seat offline or freeing it	`server/room.ts` `closed()` / `hello()`	Yes	Yes	Plausible; order depends on server timing
H3	The server answers the rejoin with a non-fatal error and the screen stays on "Connecting…" with no retry	`server/room.ts`, `src/ui/Online.tsx`	Maybe	Yes	No error is shown or logged in the test
H4	The client drops or ignores the first room state after reconnecting (version check or listener timing)	`src/net/client.ts`, `useGameController.ts`	Maybe	Yes	Code looks right, but it isn't tested under slow timing
H5	Only the test environment: an overloaded CI machine is slower than the timeout	Test scripts	Yes	Yes	Normal time is about 1 s, so 15–20 s would be extreme. Not assumed without data

Ruled out: a seat being lost when the server sleeps. Seats are stored on each connection in a way that survives the server sleeping (`serializeAttachment`).

3 · PLAN-004 · Reconnect fix (evidence first)

Player problem: a player who reloads or loses connection can be stuck on a waiting screen instead of returning to their match.

Included scope (small steps)

1. **Evidence.** When the online or quick-match test fails, it saves:
    - both players' screens, page text and browser console
    - the room messages each player received
    - the local server log, kept as a CI artifact for 7 days
2. **Reproduce.** A new stress test, `npm run test:reconnect`, repeats three kinds of cycle and times how long the game screen takes to return:
    - reload mid-match
    - a connection drop (the browser goes offline, then online)
    - quick-match pairing

    It runs normally and with the CPU slowed 4× (like a cheap phone). It can be run locally and from a manual CI workflow (DEC-011). I'll run it in batches (target 300 cycles) and record the failure rate.
3. **Isolate.** Compare failing and passing cycles to confirm or reject H1–H5.

4. **Fix the measured cause** with the smallest change, plus a regression test that fails before the fix and passes after.
    - If the cause is in seat or token handling (who owns a seat), a fresh, separate review checks the change, because that's where security matters.
    - If the cause is only test timing (H5), the fix goes in the tests, but only with stress-test data showing real rejoins work.
5. **Prove.** The stress test and CI must stay green repeatedly (acceptance below).

Excluded

- Rules, balance, `RULES_VERSION` and saves.
- Message-format changes that old app versions can't understand (if one is needed, I stop and replan).
- Publishing (that needs "PREPARE RELEASE" and your Deploy).
- Package C.
- Making tests "pass" by just waiting longer or retrying.

Roles: network/backend engineer and QA (lead), technical lead, plus security review if seat or token logic changes.

Likely files:

- `scripts/e2e-online.cjs`, `scripts/e2e-quick.cjs`, `scripts/e2e.mjs` (diagnostics, server log)
- new `scripts/e2e-reconnect.cjs`
- new `.github/workflows/reconnect-stress.yml` (manual)
- `package.json`
- then, depending on the cause: `server/room.ts` or `src/net/client.ts` or `src/ui/Online.tsx`, plus a test in `tests/room.test.ts`

Acceptance (observable)

- R1. `npm run test:reconnect` prints time-to-game for every cycle, the failures, and the 95th percentile.
- R2. The root cause is named with evidence from a failing cycle (a log or message excerpt). If it isn't reproduced after at least 300 cycles, I report **NOT REPRODUCED** with the conditions tried, and stop rather than guess.
- R3. If reproduced: the failure rate before the fix is recorded, and after the fix at least 100 slowed-down cycles have 0 failures.
- R4. A deliberately forced failure in the online test produces the diagnostics files and the server log artifact.
- R5. A regression test fails on the old code and passes on the fixed code.
- R6. `verify:full` is green on 3 consecutive CI runs; the leak check stays at 0; all unit tests stay green; no rules or save changes.

Checks: R1–R6, `verify`, the manual stress workflow, and CI.

Cost and data: RM0. GitHub Actions is free for public repositories. Diagnostics contain only test players and expire after 7 days.

Risks

- It may not reproduce outside CI. Mitigation: the slowed-down mode and the CI workflow, and a hard stop at 300 cycles with a report.
- A fix in seat handling could affect reconnect security. Mitigation: a regression test, the existing room tests, and a separate review.
- Time: likely 1–2 sessions. I'll report at the end of step 3 if the cause turns out larger than a small fix.

Rollback: revert the package commits. Players are unaffected until a release is published.

How you'll see the result: once the fix is on the preview, reload the page mid-match on your phone, or switch to airplane mode for 5 seconds and back. The match should return within a few seconds.

4 · Decisions

ID	Question	Recommendation	Status
DEC-010	How to fix the rejoin problem	Evidence and reproduction first, then fix only the measured cause (PLAN-004)	Awaiting approval
DEC-011	Run the stress test on every push?	No: on demand (manual workflow and local runs)	Awaiting approval

5 · Still open

- Pending decisions from earlier: DEC-001, DEC-004, DEC-005, and PLAN-003 (outline only).
- Your preview playtest, branch protection, your phones and your Cloudflare plan.

REP-006 · PLAN-004 done: the online "stuck on Connecting..." bug is found and fixed

Game	Fruit Market
Report date	2026-10-26 15:25 (Asia/Kuala_Lumpur)
Package	PLAN-004, approved by the owner 2026-10-26: "Approve PLAN-004 only, including its listed tests and documentation"
Build	`main @ d7e5d24` (fix in `9d9e911`)
Live game	**Unchanged** (1.1.0). The fix is on the preview; players get it only after "PREPARE RELEASE" and your Deploy
Delivery status	Verified (R1–R6 PASS)

The bug, in plain words

After a player reloaded the page or rejoined an online match, the game sometimes stayed on **"Connecting to room ..."**, even though the match was there and the server had already sent it.

Why: the screen started listening for room updates a moment *after* it first appeared. The connection was already open by then. If the server answered in that gap, the answer was missed, and the screen waited for a next message that might never come (for example, when everyone was waiting on that player).

The gap is tiny on a fast computer, so it usually worked. On a slow phone, or when the app is in the background (browsers pause drawing for hidden tabs), the gap is long and the bug happens often. The game screen had the same weakness for updates arriving just as it opened.

Fix: two small changes in the app (`src/ui/Online.tsx`, `src/ui/game/useGameController.ts`). Right after starting to listen, each screen checks what has already arrived.

- No server change, no message-format change, no rules or save change.
- Old and new app versions stay compatible.

Evidence

Check	Result
Normal timing, old code (local)	270 cycles (reloads, network drops, quick matches; normal and 4× CPU slowdown), 0 failed. The bug is too rare to see this way
Background-tab mode, **old code**	**20 of 20 reloads failed**; every one stuck on "Connecting to room" with the full game already received (`room v2 status=playing`)
Background-tab mode, **fixed code** (local)	65 cycles, 0 failed

Check	Result
Background tab plus 4× slowdown, **fixed code** (local)	100 cycles, 0 failed (reload p95 0.59 s, drop p95 0.63 s, quick match p95 1.3 s)
Same on a CI machine (where the original failures happened)	110 cycles, 0 failed (Reconnect stress run 40283316566)

Acceptance results

#	Criterion	Result	Evidence
R1	Stress test with per-cycle timings	PASS	`npm run test:reconnect` ; modes: slowdown, background tab, preview target
R2	Root cause named with evidence	PASS	Lost first update after (re)joining; see the old-code evidence above
R3	Before rate recorded; at least 100 slowed cycles with 0 failures after	PASS	20/20 failed before; 100/100 passed after (plus 175 more)
R4	Forced failure leaves evidence	PASS	Page text, room messages, console, both screenshots and the server log were saved
R5	Regression test fails before, passes after	PASS	New suite step "rejoin from a background tab"
R6	3 consecutive green CI runs; leak check 0; unit tests green; no rule or save changes	PASS	CI runs 7 (push), 8 and 9 (manual) on `d7e5d24` : all green, each with leaks 0 and 63 unit tests; no rules or save files changed

Side observations (no action taken)

- **The local test server never finishes closing a socket when the page closes it.** Cloudflare's docs say production replies automatically for this project's settings, and the real app never waits on a socket it closed itself, so no player impact was found. If it ever matters, the documented safe fix is one line in the server's close handler.
- **Playwright's offline mode doesn't break open connections**, so the stress test cuts them with a small local relay instead. A "silent" dead connection (no signal at all) isn't simulated yet. That would be part of Package C.

What changed

Added

- `npm run test:reconnect` and a manual *Reconnect stress* workflow (local or preview, optional slowdown and background-tab mode).
- Failure evidence for all online browser tests.
- The regression step in the browser suite.

Fixed (player-facing, on the preview, not yet live)

- After reloading or rejoining an online match, the game could stay stuck on "Connecting to room...", especially on slower phones or when the app was in the background.

Risks and costs

- RM0.
- The browser suite is about 1 minute longer (the regression step).
- The fix affects all online screens. Mitigation: the full suite, including the secret-price leak check, stayed green, and 545 stress cycles ran in total.
- Remaining NOT RUN: a real phone in a real background tab. Your preview playtest covers that (see below).

How to try it (preview)

1. Open https://fruit-market-staging.murshidi-hashim.workers.dev on two phones (or one phone and a computer).
2. Start an online game (Play a friend → Create room → open the link on the other device).
3. During the match, on one phone: switch to another app for 10 seconds, come back, then pull down to reload.
4. Expected: the match comes back within a few seconds both times, never stuck on "Connecting to room...".

REP-007 · PLAN-005 proposed: make 4-player games easy to find

Game	Fruit Market
Report date	2026-10-26 18:36 (Asia/Kuala_Lumpur)
Build	`main @ cc27848` (live game still 1.1.0)
Type	Planning only. **No game implementation changes were made.** One read-only test was run locally.
Approval status	PLAN-005 **awaiting approval** (DEC-012)

1 · What you noticed

You couldn't see an option for 4 human players. 4-player games are supported, but the screens don't say so.

Online room (Play a friend), as it looks today on a 360 px phone:

- The heading says **"Invite a friend"** (singular).
- There is one empty row: **"Waiting for a friend to join"**, which disappears once one person joins.
- The count "1/4" is in small grey text next to "Players".

Same phone (New game): 4 people means choosing **3** rival shops and switching each from Computer to Friend. Nothing on the screen explains that.

Quick match is always 2 strangers, by design. Nothing points groups to Play a friend instead.

2 · Check that 4 players really work (read-only, local server)

- 4 browsers joined one room by the same link. The lobby showed "Players 4/4". All 4 pressed ready, the host started, and **all 4 reached the game** (each saw 3 rivals).
- A **5th** person was turned away: "Couldn't join · This room is full."
- Limitation: tested on one machine, not 4 real phones. The automated test below makes this a permanent check.

3 · PLAN-005 · 4-player clarity

Player problem: groups of 3 or 4 think the game is for 2 and give up, or play with computers instead.

Included scope (screen text and layout only)

Screen	Today	Proposed
Online lobby, invite box	"Invite a friend"	**"Invite friends (up to 3)"**
Online lobby, player list	"Players 1/4" (small, grey), one row "Waiting for a friend to join"	**"Players 1 of 4"** in normal text, plus one row per free seat: "Open seat · send the link"
Online lobby, when full	(nothing)	"Room full · start when everyone is ready"
Home, under "Play a friend"	(nothing)	Small line: "2 to 4 players online"
Quick match screen	(nothing)	"Quick match is always 2 players. For 3 or 4, use Play a friend."
New game, under "Rival shops"	(nothing)	"Up to 4 shops. Pick Friend to pass this phone around." plus a live total, for example "4 shops: you, 2 friends, 1 computer"

Computer shops added in the lobby take seats, so the open-seat rows and the count include them. For example: 2 people + 1 computer = "Players 3 of 4", one open seat.

Also included: a new automated browser test where 4 people join one room and all reach the game, and a 5th is refused. It runs in CI on every change.

Excluded

- Rules, balance, the 4-player limit and quick-match size.
- Any server or message change.
- Renaming the buttons people tap to get there ("Play a friend", "New game"), so tests and habits keep working.
- Publishing.

Roles: UX designer (wording), gameplay engineer, QA. Self-review; low risk.

Likely files: `src/ui/Online.tsx` (lobby), `src/ui/App.tsx` (home, new game), `src/ui/styles.css`, new `scripts/e2e-four.cjs`, `scripts/e2e.mjs`.

Acceptance (observable)

- C1. The lobby shows "Players N of 4" and one open-seat row per free seat. Both update when people join or leave and when computer shops are added.
- C2. With 4 people there are no open seats and the lobby says the room is full.
- C3. The new browser test passes: 4 people join one room, all see the game, and a 5th is refused.
- C4. The new lines fit on a 360 px wide phone without overflow (screenshots of each screen in the report).
- C5. Store screenshots change only where intended (home and lobby text), checked with the visual-diff tool.
- C6. `verify:full` is green in CI; no rules, server or save changes.

Cost and data: RM0. No data changes.

Risks: longer text could wrap awkwardly on small phones; C4 checks this. Wording is a matter of taste, so you can change any line before or after.

Rollback: revert the package commit.

After it's built: open the preview, tap Play a friend → Create room, and check that the lobby clearly invites up to 3 friends.

4 · Decision

ID	Question	Recommendation	Status
DEC-012	How to make 4-player play easy to find	Screen text and open-seat rows plus a 4-person test (PLAN-005); no rule or server change	Awaiting approval

5 · Note on order

The reconnect fix (PLAN-004) is on the preview and not live yet. PLAN-005 can be built first and both can go out together in 1.1.1, which saves you one Deploy.

REP-008 · PLAN-005 done: 4-player games are easy to find

Game	Fruit Market
Report date	2026-10-26 19:05 (Asia/Kuala_Lumpur)
Package	PLAN-005, approved by the owner 2026-10-26: "Approve PLAN-005 only, including its listed tests and documentation"
Build	`main @ 04b72b5`
Live game	**Unchanged** (1.1.0). On the preview; goes live with the next release
Delivery status	Verified (C1–C6 PASS)

What changed (screen text and layout only)

Screen	Before	After
Online lobby, invite box	"Invite a friend"	"Invite friends (up to 3)". Fewer if computer shops are added, for example "(up to 2)"
Online lobby, player list	"Players 1/4" in small grey text, one row "Waiting for a friend to join"	"Players 1 of 4", plus one "Open seat · send the link" row per free seat
Online lobby, when full	Nothing	"Room full · start when everyone is ready"
Online screen	"Play a stranger..." / "send to a friend"	Quick match: "Always 2 players; for 3 or 4, start a private game below." Private game: "send to up to 3 friends... 2 to 4 players."
Home	"Play a friend"	Same button, with a small line "2 to 4 players online"
New game	Nothing explained 4 players	"Up to 4 shops. Pick Friend to pass this phone around." plus a live total, for example "4 shops: you, 2 friends, 1 computer"

Wording note: the quick-match hint says "start a private game below" instead of the planned "use Play a friend", because both options are on the same screen. The meaning is the same.

No rules, server or save changes. The buttons you tap keep their names.

Acceptance results

#	Criterion	Result	Evidence
C1	Lobby count and open seats update with joins and computer shops	PASS	New test: 1 of 4 with 3 open; 2 people + 1 computer = "3 of 4" with 1 open and "(up to 2)"

#	Criterion	Result	Evidence
C2	Full room shows no open seats and "Room full"	PASS	New test: "Players 4 of 4", 0 open, "Room full"
C3	4 people play; a 5th is refused	PASS	New test: 5th sees "This room is full"; all 4 reach the game, each seeing 3 rivals
C4	Fits a 360 px phone	PASS	Home, online and new-game screens: page exactly 360 px wide, no sideways scrolling
C5	Store screenshots change only where intended	PASS	Home 7.3% different (the button's new second line moves the rows below it). All 6 in-game screens 0.00%
C6	Full CI green; no rule, server or save changes	PASS	CI run 11 on `04b72b5` green, including the new four-player step and the preview update; diff touches only screen code, styles, tests and docs

Process note: my first `verify` run failed because I had updated a record without regenerating the views. The records check caught it, as designed (PLAN-002), and the run passed after `npm run records`.

Risks and costs

- RM0.
- The store screenshot of the home screen (`store/screenshots`) should be retaken before the next Play listing update. It still shows the old button.

How to see it

On the preview (https://fruit-market-staging.murshidi-hashim.workers.dev):

1. The home screen shows "2 to 4 players online" under Play a friend.
2. Tap Play a friend → Create room. The lobby says "Invite friends (up to 3)" and shows three open seats.
3. Tap New game → 3 rival shops → switch two to Friend. It reads "4 shops: you, 2 friends, 1 computer".

Next

Both fixes (reconnect and 4-player clarity) are on the preview. After your preview playtest, say "PREPARE RELEASE" to send them to players as 1.1.1.

REP-009 · Release 1.1.1 prepared: readiness and how to publish

Game	Fruit Market
Report date	2026-10-26 20:30 (Asia/Kuala_Lumpur)
Request	Owner: "PREPARE RELEASE"
Candidate	Version **1.1.1**: the `main` commit that adds this report (the Deploy button tags exactly that commit `v1.1.1`)
Live now	1.1.0 (`c915e49`)
Approval status	DEC-013 (publish 1.1.1) **awaiting your approval**. Nothing has been published

1 · What players get (changelog for 1.1.1)

- **Fixed:** after reloading the page or rejoining an online match, the game could stay stuck on "Connecting to room...", especially on slower phones or when the app was in the background.
- **Clearer group play:** the online lobby shows how many of the 4 seats are taken and which are still open, and the menus explain that 2 to 4 people can play together, online or by passing one phone.
- **Behind the scenes:** updates now reach players only after a preview and your approval, with automatic checks before and after.

Included packages: PLAN-001, PLAN-002, PLAN-004 and PLAN-005.

2 · Size of the change since 1.1.0

- **Game code:** 4 files, about 40 lines (`src/ui/Online.tsx` , `src/ui/App.tsx` , `src/ui/game/useGameController.ts` , `src/ui/styles.css`).
- **Server code:** unchanged. **Server settings for the live game:** unchanged (only a separate preview section was added).
- **Rules, balance, rules version (2) and saves:** unchanged. Old and new app versions can play together.
- **Version number:** 1.1.0 → 1.1.1 (shown in Settings; also the Android version name).
- **CI setting:** job time limit raised from 20 to 35 minutes (see section 3). No effect on players.

3 · Readiness gates

Gate	Check	Result	Evidence and limitation
Build ready	Clean install, records, lint, types, 63 unit tests, build	PASS	Latest green CI: run 12 on `f319548` (same game code as the candidate). Run 13 on `f472186` was **cancelled**: its browser download took 16 minutes (usually about 20 seconds), so it hit the job's 20-minute limit. The limit is now 35 minutes. The candidate's own CI run must be green before you press Deploy

Gate	Check	Result	Evidence and limitation
Play ready	Computer-played solo game; speed budget	PASS	Browser suite in CI
Play ready	**You finish a match on the preview**	**NOT RUN**	Needs you, about 10–15 minutes (owner brief, section 2)
Online ready	Two players, secret prices never leak, a player leaves	PASS	Leak check 0 on every CI run
Online ready	Rejoin and reconnect	PASS	545 stress cycles after the fix, 0 failures; regression test in CI
Online ready	4 players in one room; a 5th refused	PASS	New test in CI
Online ready	Server crash recovery	NOT RUN	Package C. Server code is unchanged since 1.1.0, so this release adds no new risk here
Release ready	Real phones (Android Chrome, iPhone Safari)	NOT RUN	Emulated phones only; Package C
Release ready	Data backup and restore	NOT RUN	No data format change in this release; Package C
Release ready	Live smoke test after deploy	Automatic	The Deploy workflow checks the live site and its rules version, then tags v1.1.1
Release ready	Rollback path	Ready	See section 5

Release-ready is not claimed, because required checks are NOT RUN. The workflow guide lets you choose between:

- **A (recommended):** do the 15-minute preview playtest first. If it passes, publish. Real phones, backups and crash recovery stay as documented gaps for Package C. That's reasonable here because 1.1.0 is already public and this release changes no server code or data.
- **B:** publish now and explicitly accept the playtest as NOT RUN (a "scoped exception"). Not recommended: it skips the only human check.

4 · How to publish (after you decide)

1. Make sure CI is green for the candidate. I'll confirm in my reply.
2. GitHub → **Actions** → **Deploy LIVE game** → **Run workflow** → type `PUBLISH` → **Run workflow**.
3. It runs the checks again, confirms the release notes for 1.1.1 exist and that `v1.1.1` isn't already used, publishes, smoke-tests the live site, and tags the release.
4. Tell me when it's green. I'll mark 1.1.1 as published (date, commit, tag) and check the live site's version.

If anything else lands on `main` before you press Deploy, ask me to re-check, because the Deploy publishes whatever `main` holds.

5 · If something goes wrong after publishing

- **The Deploy run fails:** the old version keeps running. Nothing to undo. Tell me.
- **A problem is found after a successful deploy:**
  - Fastest: Cloudflare dashboard → Workers → `fruit-market` → Deployments → roll back to the previous deployment (takes effect in seconds; no data changes in this release, so it's safe).
  - Standard: I prepare 1.1.2 with the change reverted; you press Deploy (about 30 minutes including checks).

6 · Costs and data

RM0 change. No new services. No player data or storage format changes.

7 · Decision

ID	Question	Recommendation	Status
DEC-013	Publish 1.1.1 to players?	Yes, after your preview playtest passes (option A)	Awaiting approval

To approve, reply for example: **"Playtest passed. Approve DEC-013, publish 1.1.1."** Then press Deploy. To choose option B, say so explicitly, for example "Approve DEC-013 with the playtest NOT RUN exception".

REP-010 · Theme review: does Fruit Market's theme support how it plays?

Game	Fruit Market
Report date	2026-10-26 21:47 (Asia/Kuala_Lumpur)
Request	Apply the "Game Concept and Theme Companion" (edition 1.0) to this game
Build reviewed	`main @ 83de99b` (release 1.1.1 candidate); content in `src/engine/data.ts` and `RULES.md`
Type	Creative review only. **No game, rule or content changes were made.**
Approval status	DEC-014 proposed; DEC-015 and DEC-016 **awaiting your choice**

The direction is already chosen and the game is live, so this follows the companion's path for an existing game: check what works, review coherence, propose the smallest useful test. No new concepts.

1 · Creative brief (what is confirmed, what is assumed)

Item	Status	Value
Player fantasy	Confirmed (game brief)	Run a fruit stall; outsell rival shops before your stock goes bad
Core interaction	Confirmed (rules)	Buy from the farmers, set secret prices, reveal, age overnight
Players and session	Confirmed	2–4 shops; Quick 8 / Standard 10 / Long 14 days; 10–20 min (assumption, unmeasured)
Platform	Confirmed	Phone browser and Android; touch first
Audience	**Assumption**	Casual players 12+, families sharing a phone, friends online
Desired feeling	**Assumption**	Light, cheeky tension: "will my price win?" plus "sell it before it rots"
Setting and identity	**Open**	See section 4

2 · Does the theme show up in what players do?

Repeated action	How the theme expresses it	Fit
Buying at the farmers market; poorest buys first	A morning market where small shops get first pick	**Strong**. The catch-up rule feels like market goodwill
Fruit ages each night; older fruit sells for less; spoiled fruit is binned	The freshness track and "Not so" fruit	**Strong**. This *is* the fantasy; every choice is a race against rot
Secret prices, all revealed at once; cheapest sells first	"Showdown" reveal, customer by customer	**Mixed**. Exciting and fair, but real stalls post prices openly. The reveal needs a stall-side reason (presentation idea in section 5)

Repeated action	How the theme expresses it	Fit
Customers with orders and price caps	24 named characters with quirks	**Strong** for orders and caps. The quirks are cosmetic (they never change play), which is fine but should stay labelled cosmetic
Events (Heatwave, Health Inspector, Lorry Breakdown…)	Each bends the market in a believable way	**Strong**. Heatwave aging fruit and the Inspector fining old fruit are the best theme-mechanic matches
Upgrades (Handcart, Cool Room, Juice Press, Signboard…)	Things a stall owner would actually buy	**Strong**
Favour cards	Rewards for serving a whole order yourself	**Mixed**. "Free Cash" and "Good Riddance" read as abstract game cards rather than market favours
Rival shops (animal traders)	Six characters with reaction lines	**Mixed**. Fun, but personality and play style don't match: lines are cosmetic, and how cleverly they price comes from difficulty, not character

Overall the core loop expresses the fantasy well. The weak spots are presentation: why prices are secret, a few abstract card names, and cast identity. None needs a rule change.

3 · Proposed design pillars (for your approval: DEC-015)

Pillar	Include when…	Exclude when…
1. Fresh fruit, fast money: every decision is a race against spoilage	It makes freshness visible or raises the cost of waiting	It lets fruit sit forever at no cost, or hides freshness
2. Outprice rivals in secret: commit a price blind, then see the reveal	It keeps prices hidden until the reveal and makes the reveal explain *why*	It shows rivals' prices early, or adds open haggling or chat
3. A cheerful town market: playful characters, short sessions, readable on a phone	Jokes target behaviour (a vampire who hates garlic); games stay under about 20 min	Mean or stereotype-based jokes, grind, pay-to-win, long text

4 · Coherence issues found (review before expanding)

1. **Mixed setting identity** *(owner decision, see question 1)*:
    - Customer names are Malaysian and multicultural (Aisyah, Mei Ling, Kumar, Wei Jie), and one rival is "Uncle Toad".
    - The palette is Mediterranean (terracotta, harbour teal, olive).
    - Prices are in **$**, and the fruits are blueberry, watermelon, apple and mango (two are imports in Malaysia).

    It reads as "anywhere town". That's fine, but it isn't a choice anyone made deliberately.

2. **Possible ethnic association:** "Panda Lim" pairs a Chinese surname with a panda. Most other names pair freely with professions or fantasy roles (Astronaut Aisyah, Ninja Farid), so this one stands out. Recommend renaming it (cosmetic only). "Kumar the Genie" is milder; worth a second look.

3. **Why are prices secret?** The best moment in the game has no in-world reason. That's a presentation fix, not a rule change (section 5).

4. **Repeated jokes:** each rival has 2–3 lines per situation, and a 10-day game has many reveals, so lines probably repeat. Hypothesis; check in the playtest.

5. **Icon style:** some emoji still sit beside the pixel art (🔒 Prices locked, 📊, 🙋/🤖, 🕵). Already in backlog B-33.

5 · Representative day (actual rules)

Day 4 of 10, Standard game. You and Foxy Rosa (computer). You have $14; she has $19.

1. **Morning:** the event is 🔥 **Heatwave**: watermelons sell for $4 more, but all fruit ages an extra step tonight. Tomorrow's event, shown early: 🕵 Health Inspector.
2. **Buy:** you're poorer, so you **buy first**. Fresh watermelon costs $3. You take 3 and leave the cheap blueberries: in a heatwave they'd rot by tomorrow. *Choice: profit now vs. stock that survives tonight.*
3. **Customers:** Cowboy John wants 2 watermelons and pays up to $6 + 4 = * *10** for fresh ones.
4. **Secret price:** you offer 2 at **$9**; Rosa secretly offers 2 at $8.
5. **Showdown:** prices flip; the cheapest sells first. Rosa's $8 fills John's order, and you sell nothing. Rosa quips "A dollar cheaper. Sorry, not sorry!" *Feedback:* your offer is tagged "$1 too high".
6. **Evening:** you build the ❄ **Cool Room** ($7). Your fruit now ages only every other night, and heatwaves can't touch it. *Choice: spend on protection or keep cash for tomorrow.*
7. **Night:** heatwave aging hits Rosa's shelf. Your Cool Room saves your 3 watermelons for tomorrow's customers, and tomorrow's Inspector will fine Rosa for her old fruit.

Theme-only idea for pillar 2 (not approved, no rule change): present secret pricing as *chalkboards*. Each stall chalks its price face-down, the market bell rings, and the boards flip. Same rule, with a market reason.

6 · Mini creative bible (draft: samples only, per the companion)

- **World and tone:** a sunny town market where anyone might shop: cowboys, astronauts, a vampire after sunset. Light, cheeky, never mean. Stakes: a good day's takings, never ruin.
- **Sample character:** *Grandpa Tan* (cosmetic). Silhouette: cardigan, cane, paper bag. Desire: a fair price. Flaw: remembers when apples cost ten cents. Sample reaction when you sell at full price: "Ten dollars? In my day..."
- **Sample event:** *Health Inspector*. The clipboard arrives at night; every Not-so fruit costs $2. This teaches "sell old fruit first" without a tutorial line.
- **Key screen:** the reveal. Face-down boards flip highest to lowest, the winner's coins fly to their stall, and the losing offer is tagged "$1 too high". It works without sound and without colour alone: text plus movement.
- **Do:** jokes about what characters *do*; short lines; a mix of names and roles across cultures. **Avoid:** pairing a person's ethnicity with an animal or cultural cliché; jokes longer than one line; new lore that doesn't change a decision.

7 · Riskiest creative hypothesis and the smallest test (DEC-016)

Hypothesis: first-time players understand *why* they won or lost a sale at the reveal. That is pillar 2, and it's what makes the game feel fair rather than random.

Format: a playtest with the existing game (preview or live 1.1.1). **No code needed.** It can't tell us retention, market demand or online-server behaviour.

Playtest card

Testers	3–5 people who have never played; no coaching
Setup	Their own phone. New game → Quick (8 days) → 1 computer shop (Normal)
Tasks	1) Start a game. 2) Buy fruit. 3) Set your prices. 4) After the first sales, say why you sold or didn't. 5) Finish the game
Watch for	Time to first price; whether they mention "cheapest sells first" or freshness; whether they read the event card; where they hesitate
Ask after	"What were you trying to do?" "Why did you lose that customer?" "Would you play again?" (note if they start another game unprompted)
Provisional criteria (you set them; these are suggestions, not benchmarks)	**Continue** if at least 3 of 4 can explain a lost sale by day 3. **Revise** (reveal presentation, section 5) if 2 or fewer can
Limits	Small sample; friends are kind; not proof of fun or retention

I'll record what happens (observations separate from interpretations) in the next report.

8 · Questions for you (max 3)

1. **Where is this market?**
   - (A) "Anywhere town": keep as is. *Recommended for now:* costs nothing and doesn't affect the biggest risk.
   - (B) A Malaysian pasar: RM prices, local touches; the fruit artwork could change later; much bigger art scope.
   - (C) A deliberate mix ("a market where the whole world shops"): keep the names, lean into variety.
2. **Rename Panda Lim** (and review Kumar the Genie)? Recommended: yes, a cosmetic change in a later small package.
3. **Can you find 3–5 people who've never played** for the playtest card above?

9 · Decisions

ID	Question	Recommendation	Status
DEC-014	Use the Concept and Theme Companion for Fruit Market's creative work, alongside the development workflow	Yes; one shared record system	Proposed
DEC-015	Adopt the three design pillars (section 3)?	Yes	Awaiting approval

ID	Question	Recommendation	Status
DEC-016	First creative test = the reveal-comprehension playtest (section 7)	Yes; no code	Awaiting approval

No conflict with approved decisions: release 1.1.1 (DEC-013) is unaffected, and nothing here changes rules, saves or multiplayer.

REP-011 · Setting chosen; panda rename planned (PLAN-006)

Game	Fruit Market
Report date	2026-10-26 22:16 (Asia/Kuala_Lumpur)
Owner input	Answers to REP-010's three questions
Type	Creative decisions and planning. **No game changes were made.**
Approval status	DEC-017 and DEC-018 **approved**; DEC-019 (PLAN-006) **awaiting approval**; DEC-015 and DEC-016 still awaiting approval

1 · Your answers, recorded

Question	Your answer	Recorded as
Where is this market?	"C: a deliberate 'the whole world shops here' mix"	DEC-017, approved
Rename Panda Lim?	"sure, but dont run away from the general theme coherence"	DEC-018, approved
3–5 first-time testers?	"yes"	Noted. This confirms testers are available; it isn't approval of the playtest itself (DEC-016)

2 · What option C means in practice

Written into `docs/workflow/CREATIVE-BIBLE.md` (short, one page):

- **A sunny harbour market where the whole world comes to shop.** The seaside colours you already have become the harbour; a port that trades with everywhere explains the mix of fruit.
- **Customers** keep their current mix of names and roles, freely combined, never tying a culture to a cliché.
- **Money** stays plain market dollars ($), the same for every visitor.
- **No game changes are needed for C.** It turns what's already there into a deliberate choice and sets the rules for future content.

3 · The rename, keeping the cast coherent

The rivals are animal traders with friendly nicknames: Foxy Rosa (fox), Big Benny (bear), Uncle Toad, Rocky (raccoon), Pinky (pig). The panda's lines are calm and sleepy ("Bamboo-zled you!", "very zen", "Nap time, I guess"). The new name should fit that pattern and keep the lines working.

Option	Why
Patches (recommended)	Describes the panda (black-and-white patches), like Pinky and Rocky; no real-world link
Bamboo Bea	Keeps the bamboo joke; alliterative like Big Benny

Option	Why
Sleepy Sid	Matches the napping lines

Kumar the Genie, reviewed: keep. Genie is a fantasy role like Ninja, Wizard or Astronaut; the name and role don't form a national cliché.

Checked in the code (read-only): the panda's reaction lines are found by its picture (🐼), not its name. So saved games and games in progress keep working, and they just show "Panda Lim" until they end. The panda is also one of the three daily-challenge rivals; names don't affect the score check, and PLAN-006 tests that.

4 · PLAN-006 · Rename the panda rival (cosmetic)

Scope

- Change the display name in the rival list (`src/engine/data.ts`) to your chosen name. Picture, lines, difficulty and play stay the same.
- Update the daily-challenge sentence in `RULES.md` .
- Keep "Panda Lim" reserved, so no player can take the old rival name.
- Add tests: an old save naming "Panda Lim" still gets the panda's lines; daily-challenge verification gives the same score; both names are refused as player names.

Excluded: other renames, art, lines, rules, AI or balance; renaming inside saved or in-progress games.

Acceptance

- N1. New solo games, daily challenges and online rooms show the new name.
- N2. An old save loads, still says "Panda Lim", and the panda's lines still appear.
- N3. The daily-challenge score check is unchanged for the same moves.
- N4. Both names are refused as player names.
- N5. `verify:full` is green in CI; no rule, save-format or server-logic changes.

When: built on a separate branch now and released as **1.1.2 after 1.1.1 is live**, so the 1.1.1 candidate you're testing doesn't change.

Cost: RM0. **Risk:** very low (one name plus tests). **Rollback:** revert the commit.

5 · Decisions

ID	Decision	Status
DEC-017	Setting: a market where the whole world shops	Approved (owner, 2026-10-26)
DEC-018	Rename Panda Lim, keeping the cast coherent	Approved (owner, 2026-10-26)
DEC-019	Rename scope = PLAN-006, shipped in 1.1.2	Awaiting approval
DEC-015	Three design pillars	Awaiting approval (from REP-010)
DEC-016	Reveal-comprehension playtest with your 3–5 testers	Awaiting approval (from REP-010)

REP-012 · PLAN-006 done: the panda rival is now Patches

Game	Fruit Market
Report date	2026-10-26 22:48 (Asia/Kuala_Lumpur)
Package	PLAN-006, approved by the owner 2026-10-26: "Approve PLAN-006 only, including its listed tests and documentation, name: Patches"
Build	Branch `workflow/plan-006 @ f14245a` (not on `main` yet)
Live game	**Unchanged** (1.1.0). `main` is still the 1.1.1 candidate
Delivery status	Verified (N1–N5 PASS); ships in 1.1.2, after 1.1.1

What changed

- The panda trader is now called **Patches** in new solo games, daily challenges and online rooms. Same picture (🐼), same calm, sleepy lines ("Bamboo-zled you!"), same difficulty and play.
- "Panda Lim" stays blocked as a player name, so nobody can pose as the old rival.
- The daily-challenge sentence in `RULES.md` names Patches.
- Five new tests (`tests/rename.test.ts`). Game code change: one line, plus the blocked-names list.

Acceptance results

#	Criterion	Result	Evidence
N1	New games show the new name	PASS	Rival list and daily lineup contain Patches, not Panda Lim; online rooms draw from the same list
N2	Old saves load, keep "Panda Lim", still get the panda's lines	PASS	Test saves and reloads such a game unchanged; the panda's undercut line is still chosen
N3	Daily-challenge score check unchanged	PASS	The same daily played with either name ends with identical cash and sales for every shop
N4	Both names refused for players	PASS	"Patches", "Panda Lim" and "panda lim" are reserved
N5	Full CI green; no rule, save-format or server-logic change	PASS	CI run 15 on `f14245a` green (68 unit tests plus the full browser suite)

How it reaches players

`main` holds the 1.1.1 candidate, and Deploy publishes whatever `main` holds, so this waits on its branch. After 1.1.1 is live: I merge it, you say "PREPARE RELEASE" for 1.1.2, then you press Deploy. Its player note: "The panda trader is now called Patches."

Risks and costs

- RM0, very low risk.
- **Observation (not part of this package):** the CI machines' browser download took 11–16 minutes in the last 4 runs (it used to take about 20 seconds). The 35-minute limit absorbs it, but caching the download would cut each CI run by about 10 minutes. That would be a small separate package if you want it.

REP-013 · Approvals recorded: publish 1.1.1, pillars, first playtest

Game	Fruit Market
Report date	2026-10-26 22:58 (Asia/Kuala_Lumpur)
Owner input	"yes to 1 and 2" (reply to REP-012)
Type	Decisions recorded. **No game changes.** Nothing has been published yet

1 · What you approved

ID	Decision	Status
DEC-013	Publish 1.1.1 (reconnect fix + 4-player clarity)	**Approved**. Not yet deployed: you press the button
DEC-015	The three design pillars	**Approved**. Now marked approved in the creative bible
DEC-016	First playtest with 3–5 first-time players	**Approved**. Ready to run with the kit below

One gap, recorded honestly: you didn't say whether you tried the preview yourself. It's recorded as NOT RUN. It's your call; a 5-minute look at https://fruit-market-staging.murshidi-hashim.workers.dev before deploying is still worth it (Settings should show version 1.1.1).

2 · Publish 1.1.1 now (about 3 minutes, plus about 15 minutes of automatic checks)

1. Open GitHub → **Fruit-Market** → **Actions** → **Deploy LIVE game**.
2. Tap **Run workflow**, type `PUBLISH`, then tap **Run workflow**.
3. Wait for the green tick. It checks the code, publishes, tests the live site and tags `v1.1.1`.
4. Tell me "deployed". I'll check the live version, mark 1.1.1 published, then bring in the Patches rename and prepare 1.1.2.

`main` is unchanged since the candidate (`83de99b`). These records are on a separate branch until you've deployed.

3 · Playtest kit (DEC-016)

Goal: find out whether first-time players understand *why* they won or lost a sale.

Before: use the live game after 1.1.1 is published (or the preview). Each tester uses their own phone. Don't explain the rules. You only watch and take notes.

Say only this: "This is a game about running a fruit stall. Please think aloud while you play. I can't help, but I'll ask a few questions at the end."

Tasks: New game → Quick (8 days) → 1 computer shop. Play to the end.

Note for each tester (copy this per person):

```
Tester:            Age range:        Plays phone games? (often / sometimes / rarely)
Minutes to first price set:
After day 1–3 sales, in their words: why did they sell or not sell?
Did they mention "cheapest sells first" or fresher fruit? (yes / no)
Did they read the event card? (yes / no / unsure)
Where did they hesitate or get stuck?
Finished the game? (yes / no)   Started another game without being asked? (yes / no)
"What were you trying to do?"
"Why did you lose that customer?"
"Would you play again?"
Anything that made them laugh or frown:
```

Decision rule (provisional, you set it): continue if at least 3 of 4 testers can explain a lost sale by day 3; **revise the reveal** (the chalkboard idea in REP-010) if 2 or fewer can.

Send me: the notes, as typed text or photos. I'll record what you *saw* separately from what we *think* it means, and propose the smallest next step.

4 · Records

DEC-013, DEC-015 and DEC-016 approved, with your words, in the register. Still pending from earlier: DEC-001, DEC-004, DEC-005, PLAN-003 (outline).

REP-014 · UI and UX polish: what feels unfinished, and a plan for the whole game

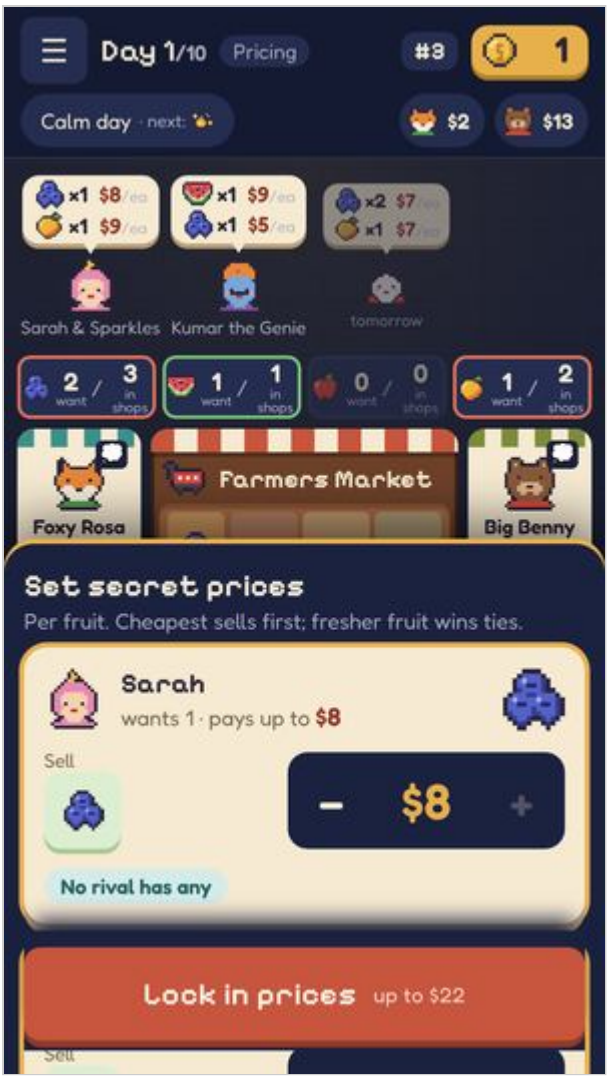
Game	Fruit Market
Report date	2026-10-26 23:06 (Asia/Kuala_Lumpur)
Owner input	"the look and feel of the game doesnt feel polished" and "the whole UI ... propose changes on the whole UI UX"
Build reviewed	1.1.1 candidate (`main @ 83de99b`), phone sizes 360×640 and 390×844
Type	Review and proposal. **No game changes were made.**
Approval status	DEC-020, PLAN-007 and PLAN-008 **awaiting approval**

1 · What I looked at

Every main screen at two common phone sizes: home, trait pick, buying, pricing, sales reveal, evening upgrades, results, trophies, online lobby and new game. The same checks the browser tests use: whether anything is cut off, overlaps or wastes space.


Day 1/10
Buying
#3
1
Calm day · next:
$10
$13
×1 $8
×1 $9
Sarah & Sparkles
walk-in
Farmers Market
Foxy Rosa
Haggler
10
Big Benny
Charmer
13
spoils
fresh
You
YOU
Big Basket
4/8
4/5
1
Your move
Tap a fruit in the market · 4/5 bought
Done buying


Day 1/10
Sales
#1
+$8
23
Calm day · next:
$11
$13
Sarah & Sparkles
Kumar the Genie
tomorrow
Foxy Rosa
Haggler
11
Big Benny
Charmer
13
SHOWDOWN 1/3
cheapest sells first
Sarah & Sparkles
×1
pays up to $8 each
"Lovely!"
You
offers 1
SOLD
$8
No rivals! 1 blueberry sold at full price.
You
YOU
Big Basket
1/8
23
Free Fruit
Skip
Selling... tap to speed up





2 · Why it doesn't feel polished: four causes

A. Things that look broken (defects, easy to fix)

1. **Sales reveal:** on shorter phones the explanation is cut off at the bottom ("...sold at full price" disappears behind the bar), and the card covers your own stall.
2. **The "+$8" money pop** flies off the top of the screen.
3. **Grammar:** "**You offers** 1" on the reveal card.
4. **Pricing panel:** the "Lock in prices" button floats over the next customer's card, which shows underneath it.
5. **Evening upgrades:** the Build buttons sit at different heights; card text is long and uneven.
6. **Trait picker:** the title wraps awkwardly and the panel covers your stall.
7. **Mixed icon styles:** emoji (🃏 favour count, ⏩ Skip, 💭, 📊, 🔒, 🧑‍🤝‍🧑/🤖) beside the pixel art.
8. **Placeholders that look unfinished:** the grey "walk-in" and "tomorrow" ghosts and the dashed "?" bubble.

9. **Trophies:** achievements have no icons, and locked and unlocked look the same (all grey text).

B. Too much on screen at once (layout)

- Before the table you see: the top bar, a second row (event, rival cash), customers, and a 4-box demand strip ("1+ want / 3 in shops"). That's four bands of small information.
- **Numbers shown twice:** your cash (top right and your stall) and rivals' cash (top chips and their stalls).
- **Rival stalls** take about 40% of the width to show mostly "0" rows, which squeezes the farmers market and makes its price labels ($1 $2 $3) tiny.
- **Your stall** is a big grid of empty pastel boxes with tiny "−$2" labels. Hard to read at a glance.

C. No single focus per moment Every phase shows the whole table at the same size. When you're pricing, the customers matter most; when you're buying, the market does. Right now nothing steps forward, so your eye has nowhere to land.

D. The visual system isn't consistent

- Two fonts are used without a clear rule: the pixel font appears in titles, numbers and small labels alike.
- Cards come in several styles: cream paper, navy chips, outlined boxes.
- Colours carry meaning (red outline, green outline, teal), but it isn't explained or applied consistently.
- Feedback and motion are uneven: some actions animate, others just switch.



3 · Proposal: three phases, smallest risk first

Phase 1 · Fix what looks broken (PLAN-007, build now, low risk)

All nine items in **A**, presentation only:

- the reveal card fits every phone and stops covering your stall
- the money pop stays on screen
- the grammar and other copy are corrected
- the pricing button sits in a solid footer
- upgrade cards are aligned and their text shortened
- the trait picker fits
- the remaining emoji become pixel icons (backlog B-33)
- the walk-in and tomorrow placeholders get proper art
- trophies get icons and a clear locked look

Ships with the Patches rename as **1.1.2**.

Phase 2 · Calm the table (PLAN-008 mockups first, then a build plan)

- **One compact top bar:** menu · day · phase · today's event icon · your cash and rank. Rivals' cash appears only on their stalls.
- **Slim rival stalls:** face, name, cash, and only the fruit they actually hold. This frees width so the farmers market and its prices get bigger.
- **Plain-language demand:** replace "1+ want / 3 in shops" with icons and words on the customer bubbles, for example "2 shops selling: price war".
- **Your stall:** show your fruit as items with "nights left", not a grid of empty boxes.
- **One focus per phase:** buying enlarges the market; pricing brings customers forward and dims the rest; the reveal gets the full stage; the evening shows only the upgrades.

Phase 3 · Make it feel finished (after Phase 2 is approved)

- **A small design system:** one type scale (pixel font for titles and big numbers only), one card style, spacing on a fixed grid, and colour roles (green good, red risk, gold money, teal info).
- **Motion and feedback:** phase transitions, a press effect on every button, and the **chalkboard reveal** from REP-010 (price boards flip). The reveal waits for the playtest result, because it's pillar 2.
- **Sound paired with the key moments** (backlog B-34), plus the harbour-market identity (DEC-017) on the home screen: a harbour backdrop, crates, chalkboards.

4 · Order of work (recommended)

1. **Now:** publish 1.1.1 and run the first playtest (DEC-016) on the *current* look. That gives a baseline, so we can show whether the new UI really makes the game clearer.
2. **Now, in parallel:** build Phase 1 (PLAN-007). Low risk; goes out as 1.1.2.
3. **Then:** PLAN-008, interactive mockups of the buy, price, reveal and evening screens plus home, as a page you can open on your phone, with two options where there's a real choice (for example, the rival stalls). **No game code.** You pick.
4. **Then:** a build plan for the chosen design (Phases 2 and 3), released as **1.2.0**, followed by a second playtest compared with the baseline.

5 · PLAN-007 acceptance (Phase 1)

- P1. Every item in **A** has before-and-after screenshots at 360×640 and 390×844.
- P2. Nothing is cut off or overlapping, and nothing scrolls sideways, on those sizes (automated check on each phase screen).
- P3. The visual-diff tool shows changes only on the screens being fixed.
- P4. The speed budget still passes (no long tasks over 200 ms; page size within budget).
- P5. `verify:full` is green; no rule, save or server changes.

PLAN-008 acceptance (mockups): an interactive page you can open on your phone, covering 5 screens, with clear options where a choice is needed, and every idea mapped to a problem above. No game code.

6 · Costs and risks

- RM0. No new tools or services.
- **Risk:** Phase 2 changes how the table is laid out, which could confuse returning players. Mitigation: mockups first, then the second playtest.
- **Risk:** polish can expand forever. Mitigation: each phase has a fixed list and its own approval.

7 · Questions for you (max 3)

1. **What bothers you most?**

   - (a) cluttered screens
   - (b) things that look broken
   - (c) it lacks animation or "feel"
   - (d) the art style itself

   (This sets the Phase 2–3 priorities. If it's (d), that's a bigger art decision.)

2. **Is there a game whose look you'd call "polished"?** A name or a screenshot helps me aim the mockups.

3. **Approve PLAN-007 (Phase 1 fixes) and PLAN-008 (mockups)?**

8 · Decisions

ID	Question	Recommendation	Status
DEC-020	How to polish the UI and UX	Three phases; fix defects first; mockups before any layout change; a playtest before and after	Awaiting approval
PLAN-007	Phase 1: fix what looks broken	Build now; ship in 1.1.2	Awaiting approval
PLAN-008	Interactive mockups for Phases 2–3	Design only, no game code	Awaiting approval

REP-015 · Art direction: what "It Takes Two" can teach a 2D phone game

Game	Fruit Market
Report date	2026-10-26 23:18 (Asia/Kuala_Lumpur)
Owner input	Bothers most: (a) clutter, (c) lack of animation and feel, (d) the art style itself. Polished reference: *It Takes Two*
Type	Creative direction options and a revised plan. **No game changes.**
Approval status	DEC-021 and the revised PLAN-008 **awaiting approval**; PLAN-007 still awaiting approval

1 · What this changes

REP-014 started with fixing broken-looking details. Your answer puts the priority elsewhere: **the art style, the feel, and the clutter**. That's a bigger and more important decision, so it should be designed and tested before anything is built.

2 · What makes *It Takes Two* feel polished (the qualities we can borrow)

It Takes Two (Hazelight, 2021) is a 3D console game, so we can't copy its look or technology, and we must not reuse any of its art or characters. Its polish comes from qualities that translate to 2D:

Quality	In *It Takes Two*	What it could mean for Fruit Market
Toy-like, touchable materials	The world is shrunk to toy size: wood, felt, clay, fabric	Stalls of painted wood, fruit like shiny clay or wooden toys, cloth awnings, chalkboard price tags
Chunky, rounded shapes	Soft silhouettes, big heads, readable at a glance	Rounder cards and buttons, bigger and simpler characters, fewer thin lines
Warm light and depth	Soft shadows, glow, layers of foreground and background	Gentle drop shadows, a warm late-afternoon harbour backdrop, depth between the table and the cards
Everything reacts	Squash and stretch, bounces, particles on every action	Fruit hops into your basket, coins bounce, buttons squish, the price reveal *flips*
Expressive characters	Faces and poses show feeling all the time	Customers smile, frown or shrug at the reveal; rivals react with a pose, not just a text line
One thing to look at	The camera always frames the important moment	Each phase zooms into what matters (market, customers, reveal) and dims the rest

3 · Three art directions (choose one)

	A · Polished pixel	**B · Toy-box tabletop (recommended)**	**C · 3D toy diorama**
Idea	Keep pixel art; make it consistent and animated	A smooth 2D illustrated style: toy-like, rounded, warm, with soft shadows and textures	Real 3D: a miniature market you look down on

	A · Polished pixel	B · Toy-box tabletop (recommended)	C · 3D toy diorama
Fixes your (d) art style?	**Weakly**: same look, just cleaner	**Yes**: a clearly new, warmer, toy-like look	Yes, and closest to the reference
Fixes (c) feel	Partly	Yes: built around motion from the start	Yes
Phone speed	Excellent	Good (vector art; we keep the speed budget)	Risky on cheap phones; heavier download
Scope	Small	Medium-large: redraw fruits, ~30 characters, rivals, ~40 icons, UI	Very large: a new rendering engine and 3D models
Who makes the art	Claude, in code (as now)	Claude, in code as vector art (consistent, revisable), **or** a hired illustrator later	3D modelling is beyond what Claude can do well; it needs an artist (paid)
Cost	RM0	RM0 if Claude draws it; an illustrator would be a separate paid decision	Paid artist plus large engineering
Risk	Doesn't satisfy (d)	Art quality by Claude is good but not AAA; judge it on the style test before committing	Highest cost, slowest; could stall the game

Recommendation: B, proven by a small style test first. It addresses all three of your concerns, stays fast on phones, costs nothing if I draw it, and can be judged before any real work. If the test looks too plain, the next step would be to brief a paid illustrator. That would be a separate, priced decision under your RM500/month ceiling.

4 · Revised PLAN-008 · Style test plus layout mockups (design only, no game code)

1. **Style test:** the *same* buying screen drawn three ways (A, B, C as a flat render), plus one fruit at its four freshness stages, one customer and one rival in direction B. It's a page you open on your phone.
2. **Layout mockups in the chosen direction** for buy, price, reveal, evening and home, with:
   - the clutter fixes from REP-014 (one top bar, slim rival stalls, demand in plain words)
   - one focus per phase
   - motion notes (what moves, when) and a short animated example of the reveal flip
3. **Your choice:** you pick the direction and the layout options. Only then do I propose the build plan(s).

Acceptance: M1, a phone-viewable page with the style test (3 directions on one screen) and the 5 mockup screens. M2, every idea is linked to a problem in REP-014 or a quality in section 2. M3, no game code changed. All art is original; reference games are inspiration only.

5 · Where PLAN-007 (the 9 broken-looking fixes) fits now

You didn't pick (b), and a new art style would redraw most of those screens anyway. Two options:

- **Keep it small (recommended):** do only the fixes that survive any art change. That's the grammar ("You offers"), the money pop going off-screen, and the sales text being cut off. Ship them in 1.1.2 with Patches; park the rest until the new style.
- Or drop PLAN-007 entirely and fold everything into the new design.

6 · Order of work

1. Publish 1.1.1 (your Deploy) and run the first playtest on the current look (baseline).
2. PLAN-008 style test, then you choose a direction.
3. Layout mockups in that direction, then you choose.
4. A build plan for the new look and feel (likely several packages), released as **1.2.0**, then a second playtest.

7 · Questions (max 3)

1. **Approve the revised PLAN-008** (style test plus mockups, no game code)?
2. **PLAN-007:** keep it small (3 fixes in 1.1.2) or drop it?
3. **Is a paid illustrator possible later,** if the B style test isn't good enough? A yes or no now helps me plan; any spending would still need its own quote and approval.

8 · Decisions

ID	Question	Recommendation	Status
DEC-021	Art direction for the polish	Explore A/B/C in a style test; recommend B (toy-box tabletop)	Awaiting approval
PLAN-008 (revised)	Style test plus layout mockups, design only	Approve	Awaiting approval
PLAN-007	Broken-looking fixes	Reduce to 3 fixes that survive any art change	Awaiting your choice